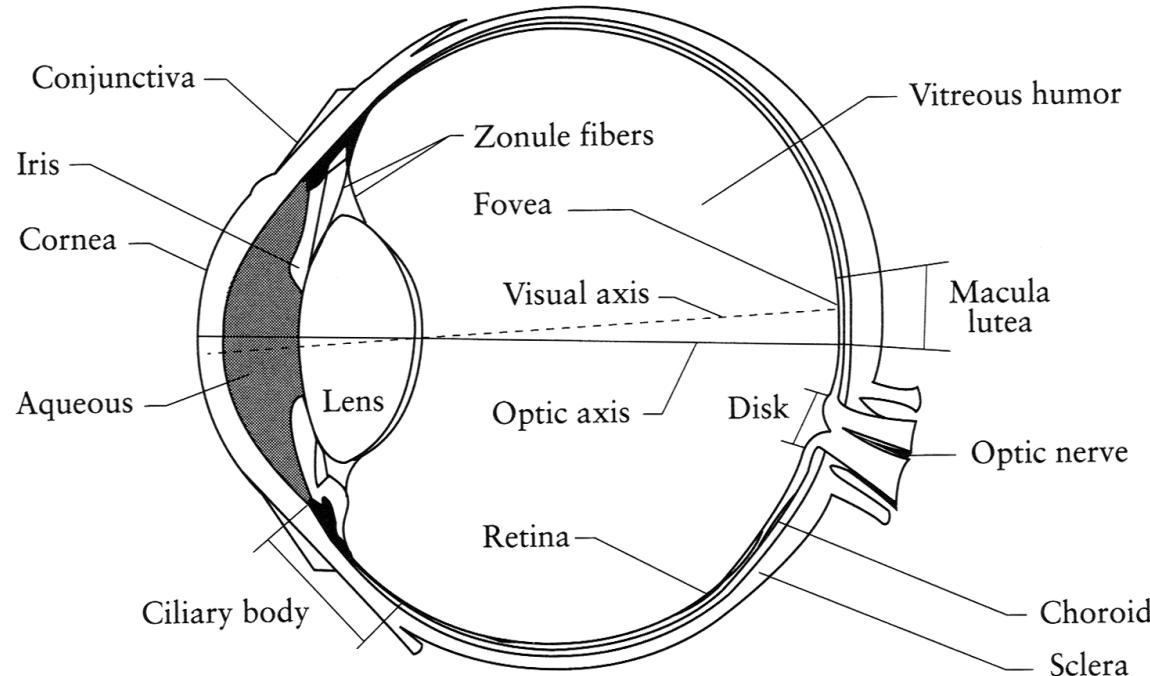


# Lecture 3 Color & Color Quantization

Multimedia Systems  
Spring 2020

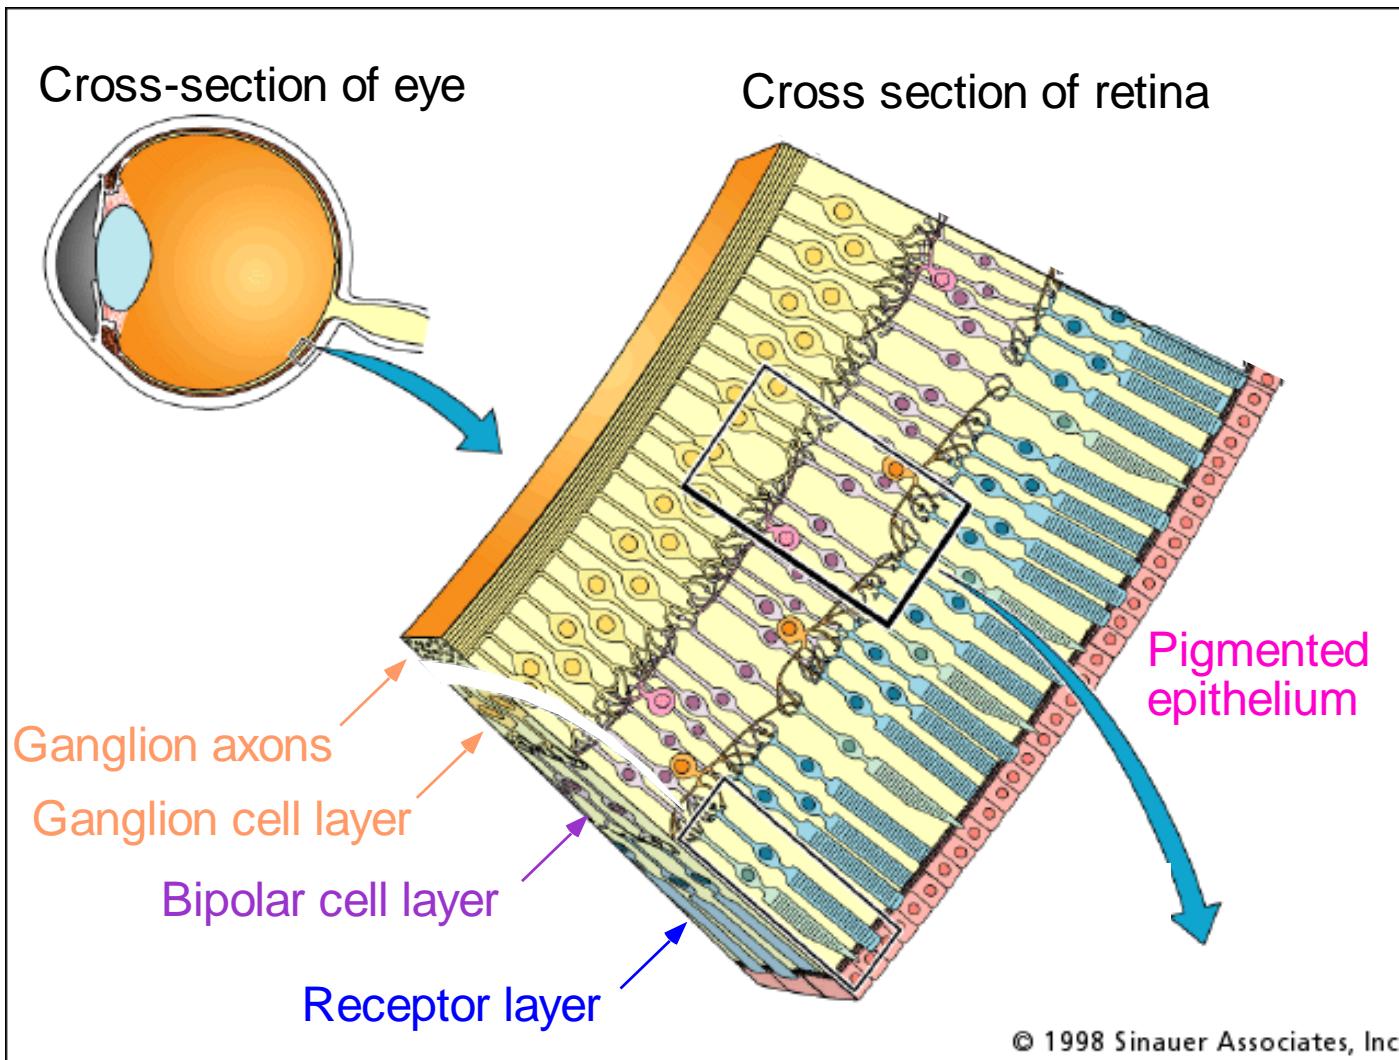
# The Eye

---

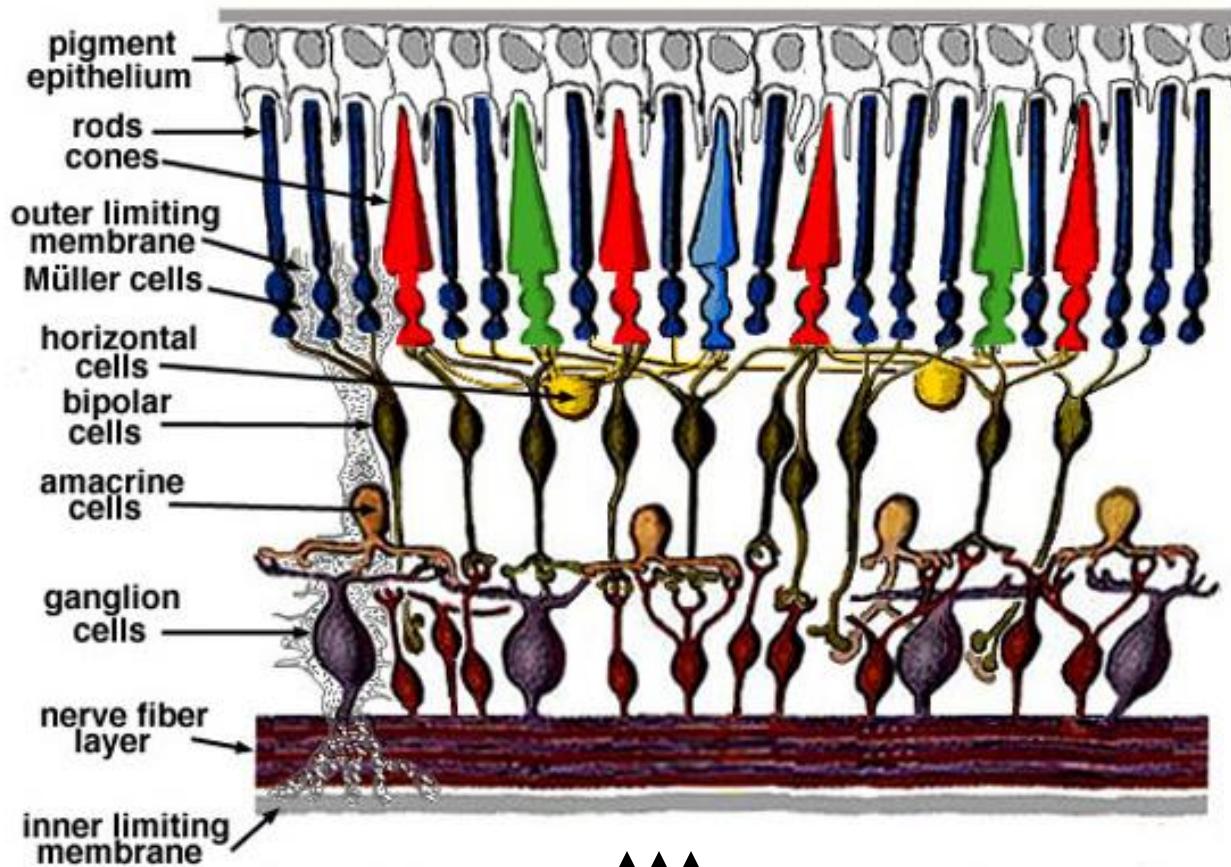


- ▶ The human eye is a camera!
  - Iris – colored annulus with radial muscles
  - Pupil – the hole (aperture) whose size is controlled by the iris
  - What's the “film”?
    - photoreceptor cells (rods and cones) in the retina

# The Retina

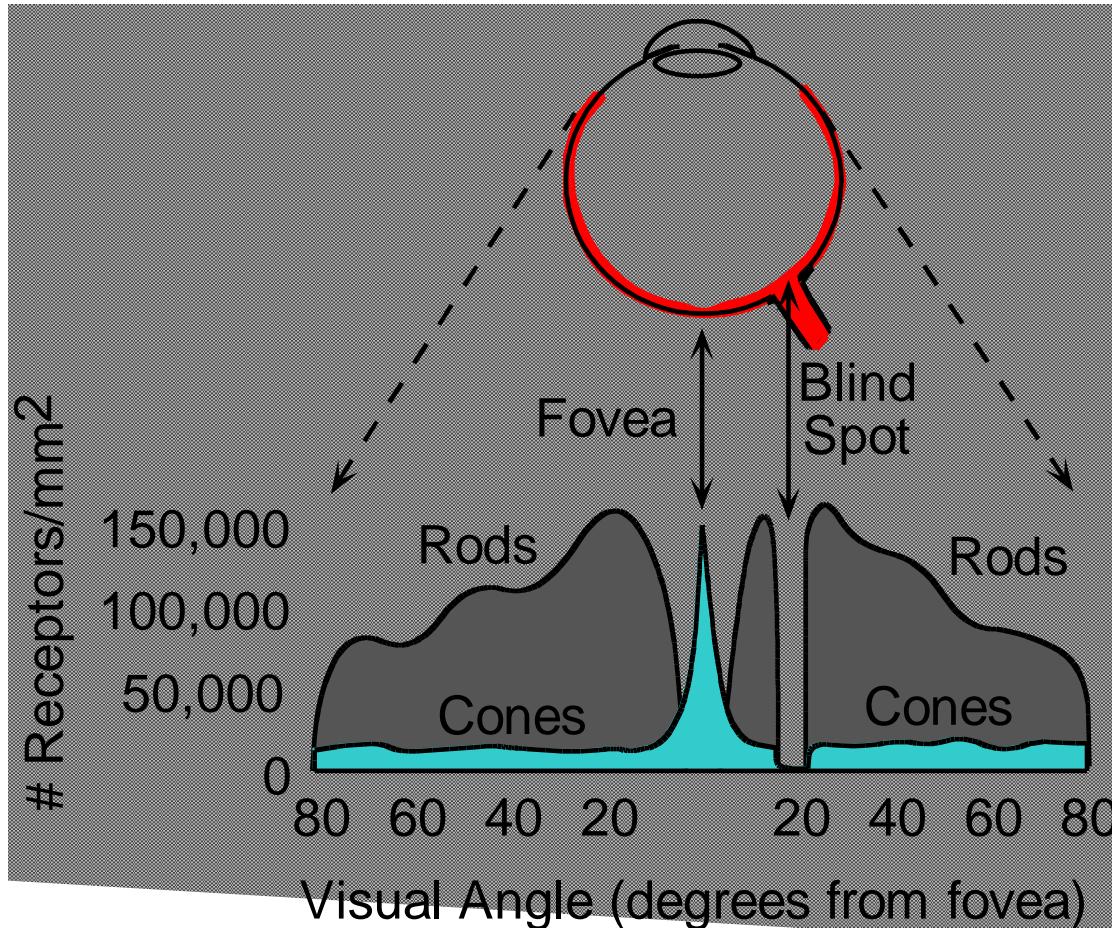


# Retina up-close



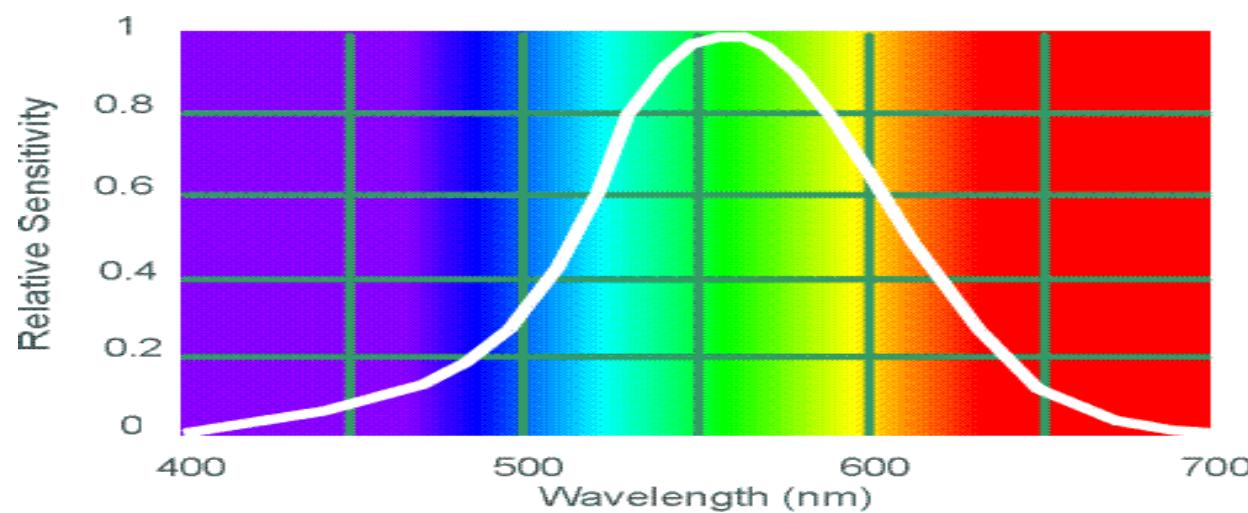
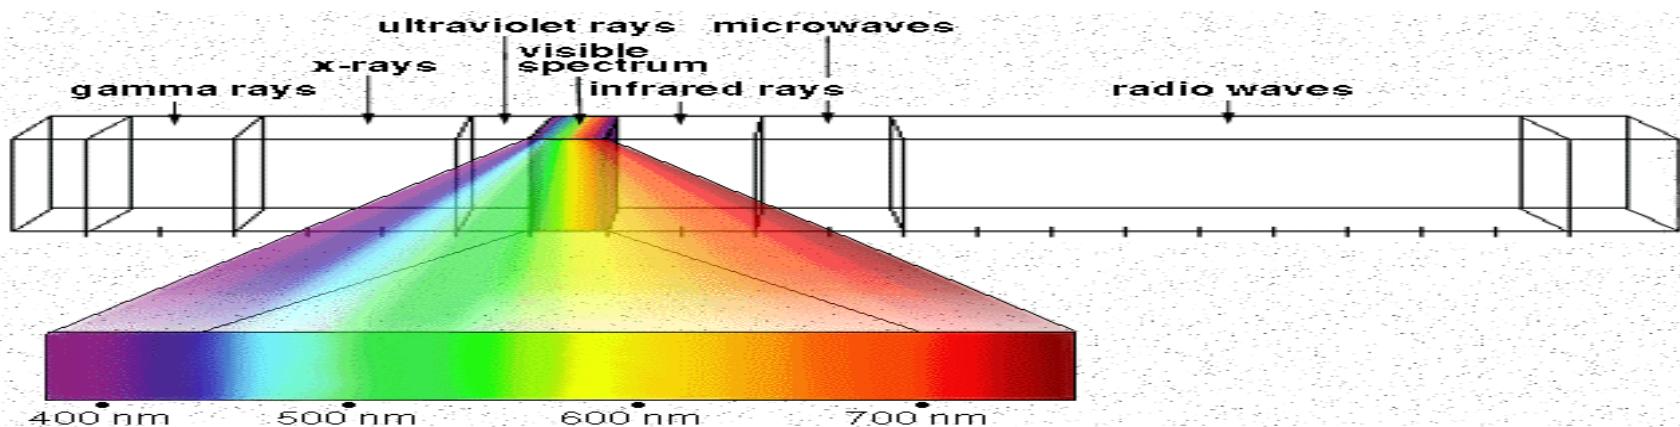
Light

# Distribution of Rods and Cones



Night Sky: why are there more stars, off-center?

# Electromagnetic Spectrum

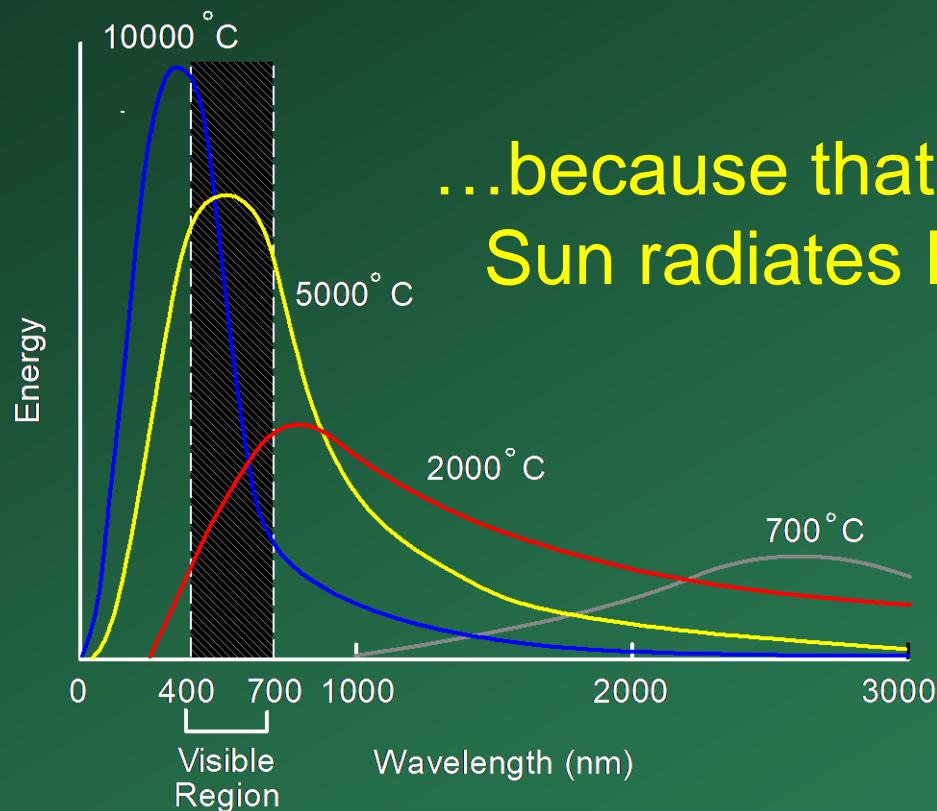


Human Luminance Sensitivity Function

# Visible Light

Plank's law for Blackbody radiation  
Surface of the sun: ~5800K

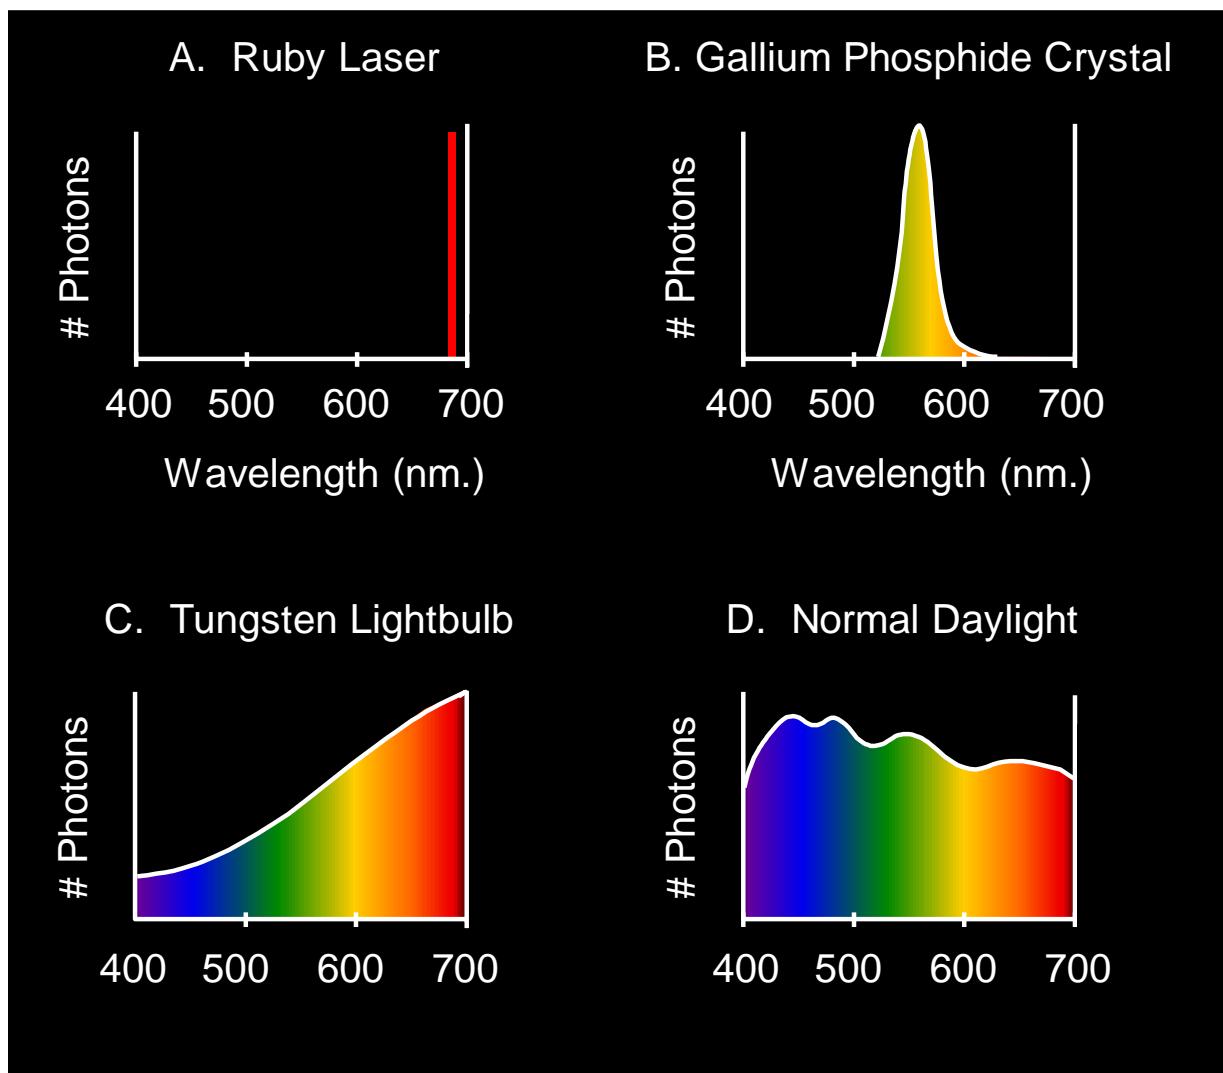
Why do we see light of these wavelengths?



...because that's where the Sun radiates EM energy

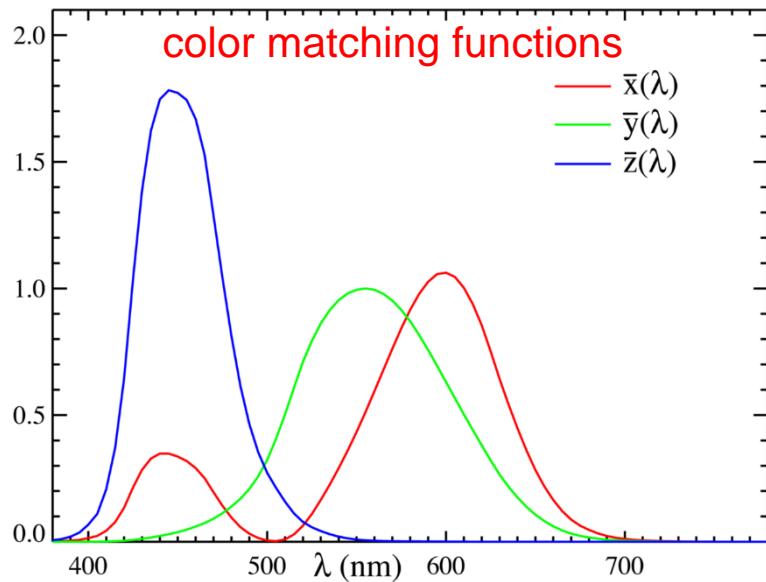
# The Physics of Light

---



# CIE Color Space

- ▶ CIE (Commission Internationale d'Eclairage) came up with 3 hypothetical lights X, Y, and Z with these spectra:



spectral power distribution of the illuminant

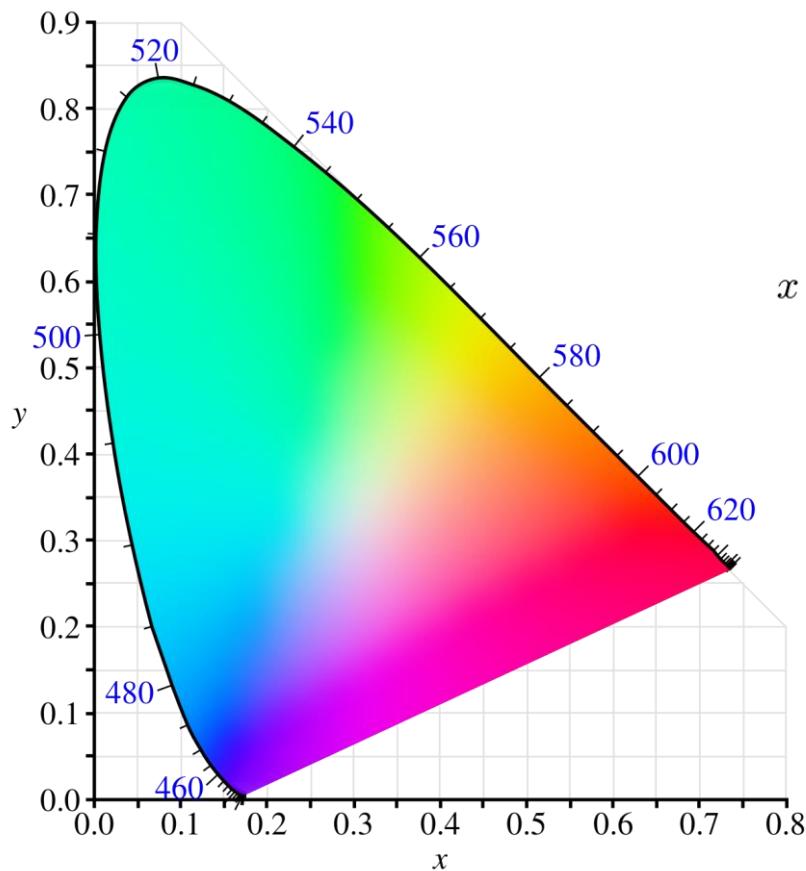
$$X = \int_0^{\infty} I(\lambda) \bar{x}(\lambda) d\lambda$$

$$Y = \int_0^{\infty} I(\lambda) \bar{y}(\lambda) d\lambda$$

$$Z = \int_0^{\infty} I(\lambda) \bar{z}(\lambda) d\lambda$$

- ▶ Idea: any wavelength  $\lambda$  can be matched perceptually by *positive* combinations of X,Y,Z
  - CIE created table of XYZ values for all *visible* colors

# CIE Chromaticity Diagram (1931)



- ▶ For simplicity, we often project to the 2D plane
- ▶ Also normalize

$$x + y + z = 1$$

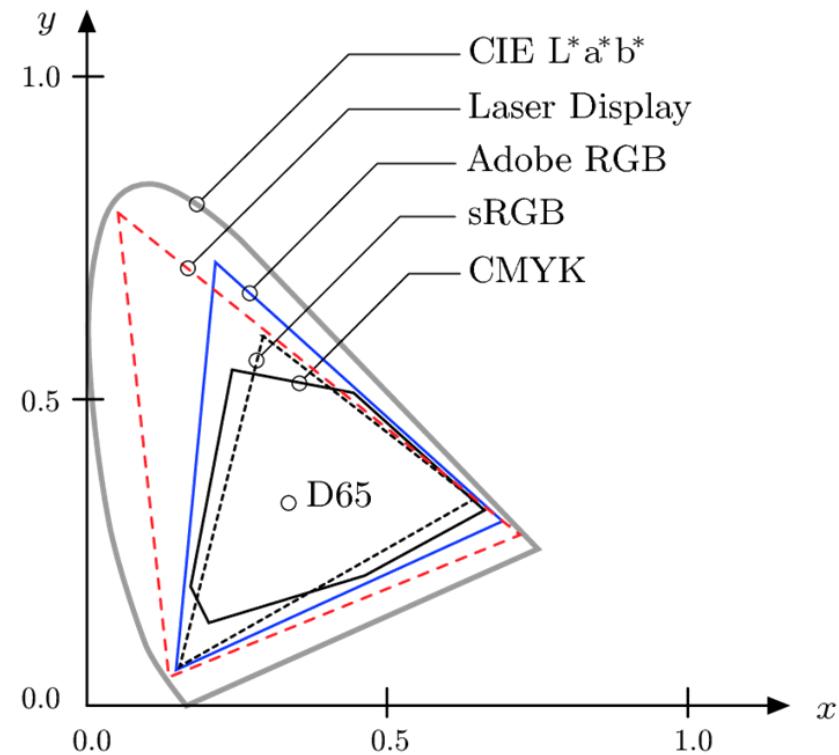
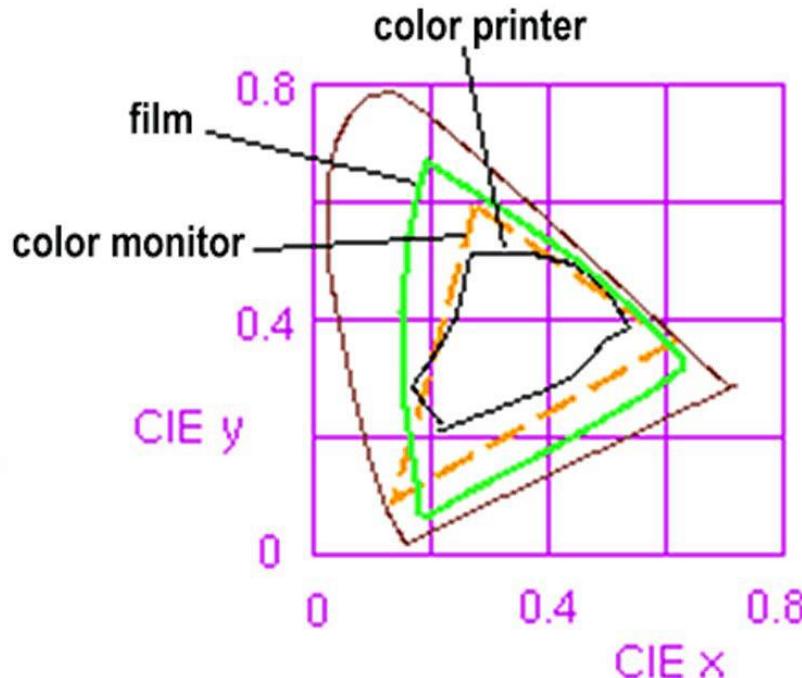
$$x = \frac{X}{X + Y + Z} \quad y = \frac{Y}{X + Y + Z} \quad z = \frac{Z}{X + Y + Z}$$

- ▶ Note: Inside horseshoe visible, outside invisible to eye

**Note:** Look up  $x, y$   
**Calculate**  $z$  as  $1 - x - y$

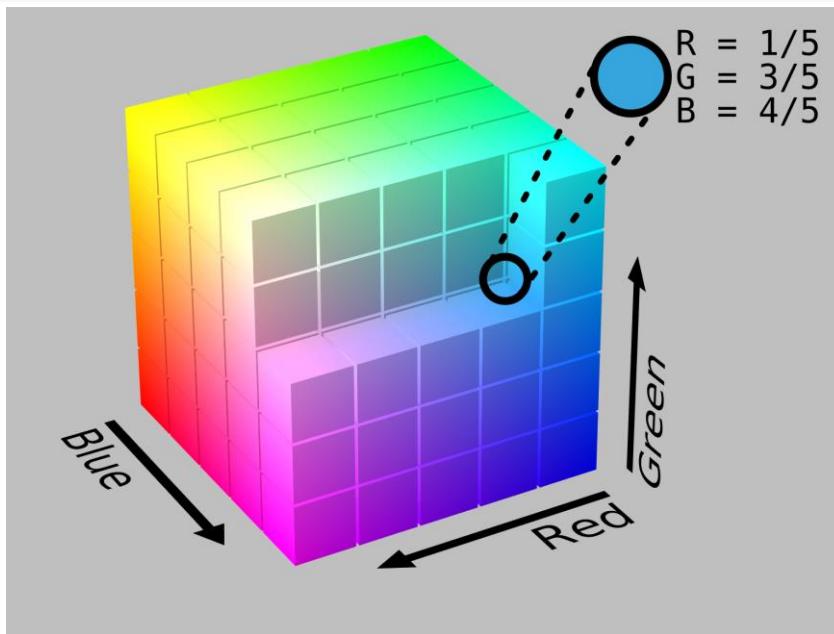
# Device Color Gamuts

- ▶ The RGB color cube sits within CIE color space
- ▶ We can use the CIE chromaticity diagram to compare the gamuts of various devices
- ▶ E.g. compare color printer and monitor color gamuts



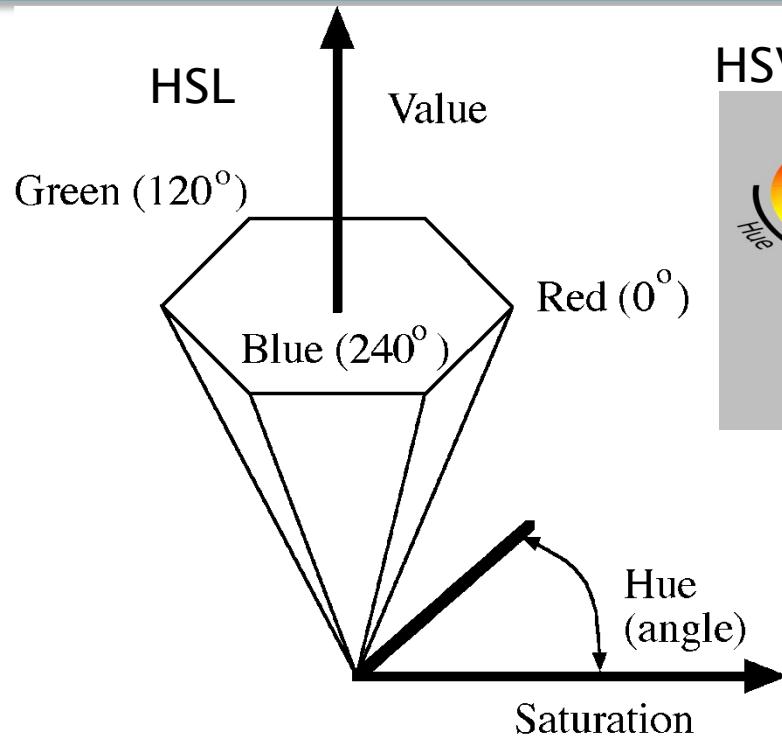
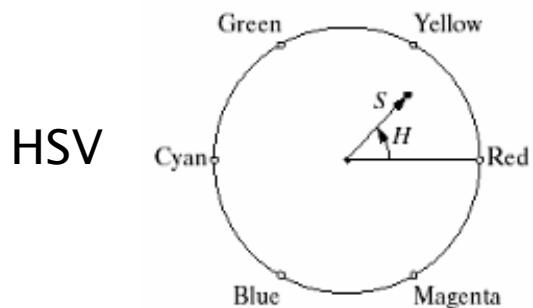
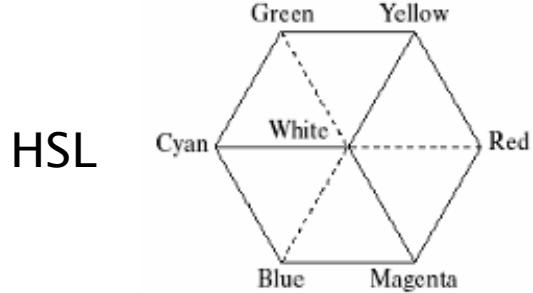
# RGB color space

---



- ▶ RGB cube
  - Easy for devices
  - But not perceptual
  - Where do the grays live?
  - Where is hue and saturation?

# HSV



- ▶ Hue, Saturation, Value (Intensity)
  - RGB cube on its vertex
- ▶ Decouples the three components (a bit)
- ▶ Use `rgb2HSV()` and `HSV2RGB()` in Matlab

# RGB to HSV Conversion

---

- ▶ Define the following values

$$C_{\text{high}} = \max(R, G, B), \quad C_{\text{low}} = \min(R, G, B), \quad \text{and}$$
$$C_{\text{rng}} = C_{\text{high}} - C_{\text{low}}.$$

- ▶ Find **saturation** of RGB color components ( $C_{\text{max}}=255$ )

$$S_{\text{HSV}} = \begin{cases} \frac{C_{\text{rng}}}{C_{\text{high}}} & \text{for } C_{\text{high}} > 0 \\ 0 & \text{otherwise} \end{cases}$$

- ▶ And **luminance (Value)** value

$$V_{\text{HSV}} = \frac{C_{\text{high}}}{C_{\text{max}}}$$

# RGB to HSV Conversion

---

- ▶ Normalize each component using

$$R' = \frac{C_{\text{high}} - R}{C_{\text{rng}}} \quad G' = \frac{C_{\text{high}} - G}{C_{\text{rng}}} \quad B' = \frac{C_{\text{high}} - B}{C_{\text{rng}}}$$

- ▶ Calculate preliminary hue value  $H'$  as

$$H' = \begin{cases} B' - G' & \text{if } R = C_{\text{high}} \\ R' - B' + 2 & \text{if } G = C_{\text{high}} \\ G' - R' + 4 & \text{if } B = C_{\text{high}} \end{cases}$$

- ▶ Finally, obtain final **hue** value by normalizing to interval  $[0,1]$

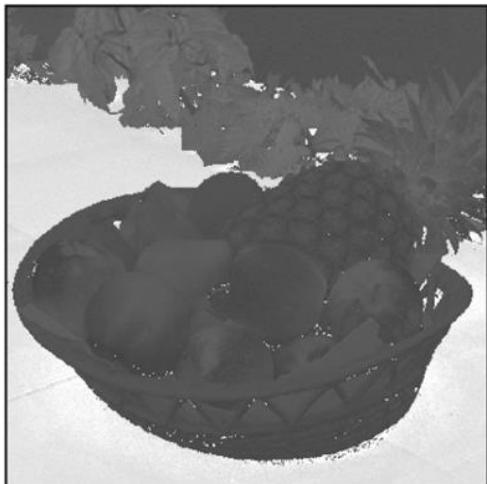
$$H_{\text{HSV}} = \frac{1}{6} \cdot \begin{cases} (H' + 6) & \text{for } H' < 0 \\ H' & \text{otherwise} \end{cases}$$

# RGB to HSV Conversion

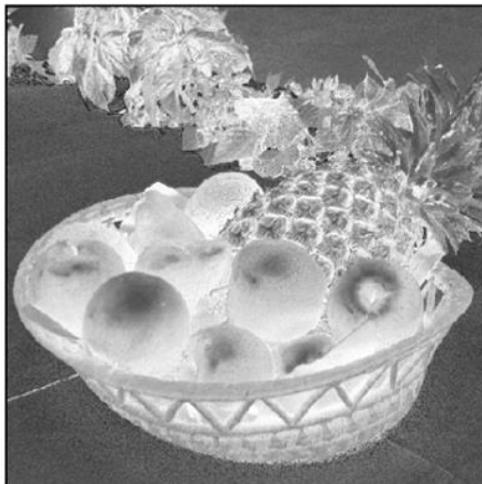
---



Original RGB image



$H_{HSV}$



$S_{HSV}$



$V_{HSV}$

# TV Color Spaces – YUV, YIQ, $YC_bC_r$

---

- ▶ YUV, YIQ: color encoding for analog NTSC and PAL
- ▶  $YC_bC_r$ : Digital TV encoding
- ▶ Key common ideas:
  - Separate luminance component Y from 2 chroma components
  - Instead of encoding colors, encode color differences between components (maintains compatibility with black and white TV)

# YUV

---

- ▶ Basis for color encoding in analog TV in north america (NTSC) and Europe (PAL)
- ▶ Y components computed from RGB components as

$$Y = 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B$$

- ▶ UV components computed as:

$$U = 0.492 \cdot (B - Y) \quad \text{and} \quad V = 0.877 \cdot (R - Y)$$

# YUV

---

- ▶ Entire transformation from RGB to YUV

$$\begin{pmatrix} Y \\ U \\ V \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ -0.147 & -0.289 & 0.436 \\ 0.615 & -0.515 & -0.100 \end{pmatrix} \cdot \begin{pmatrix} R \\ G \\ B \end{pmatrix}$$

- ▶ Invert matrix above to transform from YUV back to RGB

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} = \begin{pmatrix} 1.000 & 0.000 & 1.140 \\ 1.000 & -0.395 & -0.581 \\ 1.000 & 2.032 & 0.000 \end{pmatrix} \cdot \begin{pmatrix} Y \\ U \\ V \end{pmatrix}$$

# YIQ

---

- ▶ Original NTSC used variant of YUV called YIQ
- ▶  $Y$  component is same as in YUV
- ▶ Both U and V color vectors rotated and mirrored so that

$$\begin{pmatrix} I \\ Q \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \cdot \begin{pmatrix} \cos \beta & \sin \beta \\ -\sin \beta & \cos \beta \end{pmatrix} \cdot \begin{pmatrix} U \\ V \end{pmatrix}$$

**2D rotation matrix**

where  $\beta = 0.576$  (33 degrees)

# YC<sub>b</sub>C<sub>r</sub>

---

- ▶ Internationally standardized variant of YUV
- ▶ Used for digital TV and image compression (e.g. JPEG)
- ▶  $Y, C_b, C_r$  components calculated as

$$Y = w_R \cdot R + (1 - w_B - w_R) \cdot G + w_B \cdot B$$

$$C_b = \frac{0.5}{1 - w_B} \cdot (B - Y)$$

$$C_r = \frac{0.5}{1 - w_R} \cdot (R - Y)$$

- ▶ Inverse transform from YC<sub>b</sub>C<sub>r</sub> to RGB

$$R = Y + \frac{1 - w_R}{0.5} \cdot C_r$$

$$G = Y - \frac{w_B \cdot (1 - w_B) \cdot C_b - w_R \cdot (1 - w_R) \cdot C_r}{0.5 \cdot (1 - w_B - w_R)}$$

$$B = Y + \frac{1 - w_B}{0.5} \cdot C_b$$

# YC<sub>b</sub>C<sub>r</sub>

---

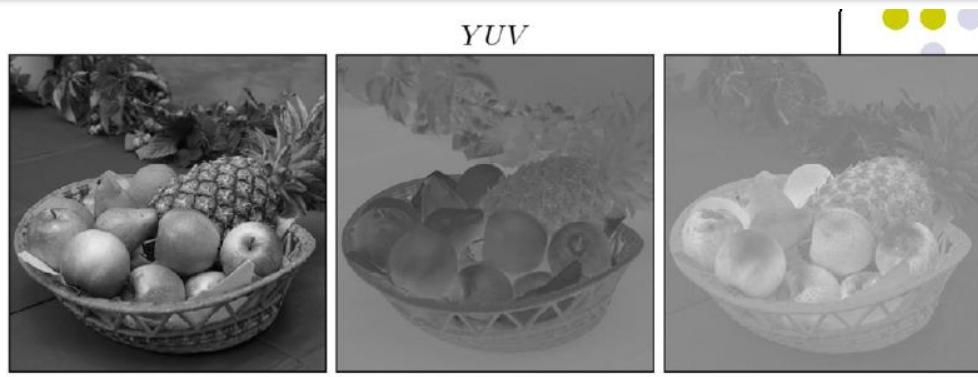
- ▶ ITU recommendation BT.601 specifies values:  
 $w_R = 0.299, w_B = 0.114, w_G = 1 - w_B - w_R = 0.587$
- ▶ Thus the transformation

$$\begin{pmatrix} Y \\ C_b \\ C_r \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ -0.169 & -0.331 & 0.500 \\ 0.500 & -0.419 & -0.081 \end{pmatrix} \cdot \begin{pmatrix} R \\ G \\ B \end{pmatrix}$$

- ▶ And the inverse transformation becomes

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} = \begin{pmatrix} 1.000 & 0.000 & 1.403 \\ 1.000 & -0.344 & -0.714 \\ 1.000 & 1.773 & 0.000 \end{pmatrix} \cdot \begin{pmatrix} Y \\ C_b \\ C_r \end{pmatrix}$$

# Comparing YUV, YIQ and $YC_bC_r$

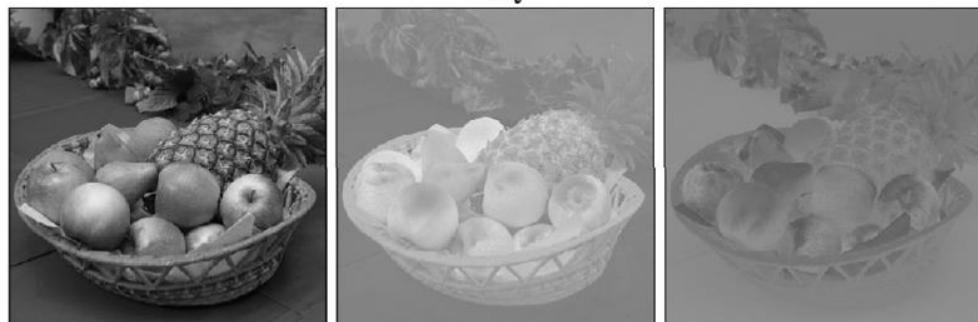


$Y$

$U$

$V$

$YIQ$

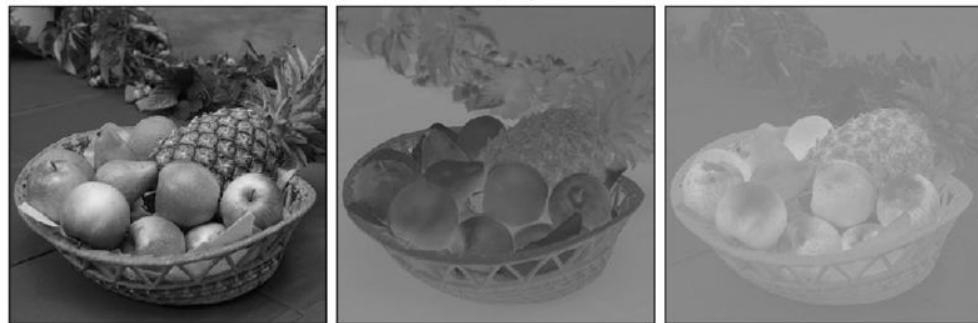


$Y$

$I$

$Q$

$YC_bC_r$



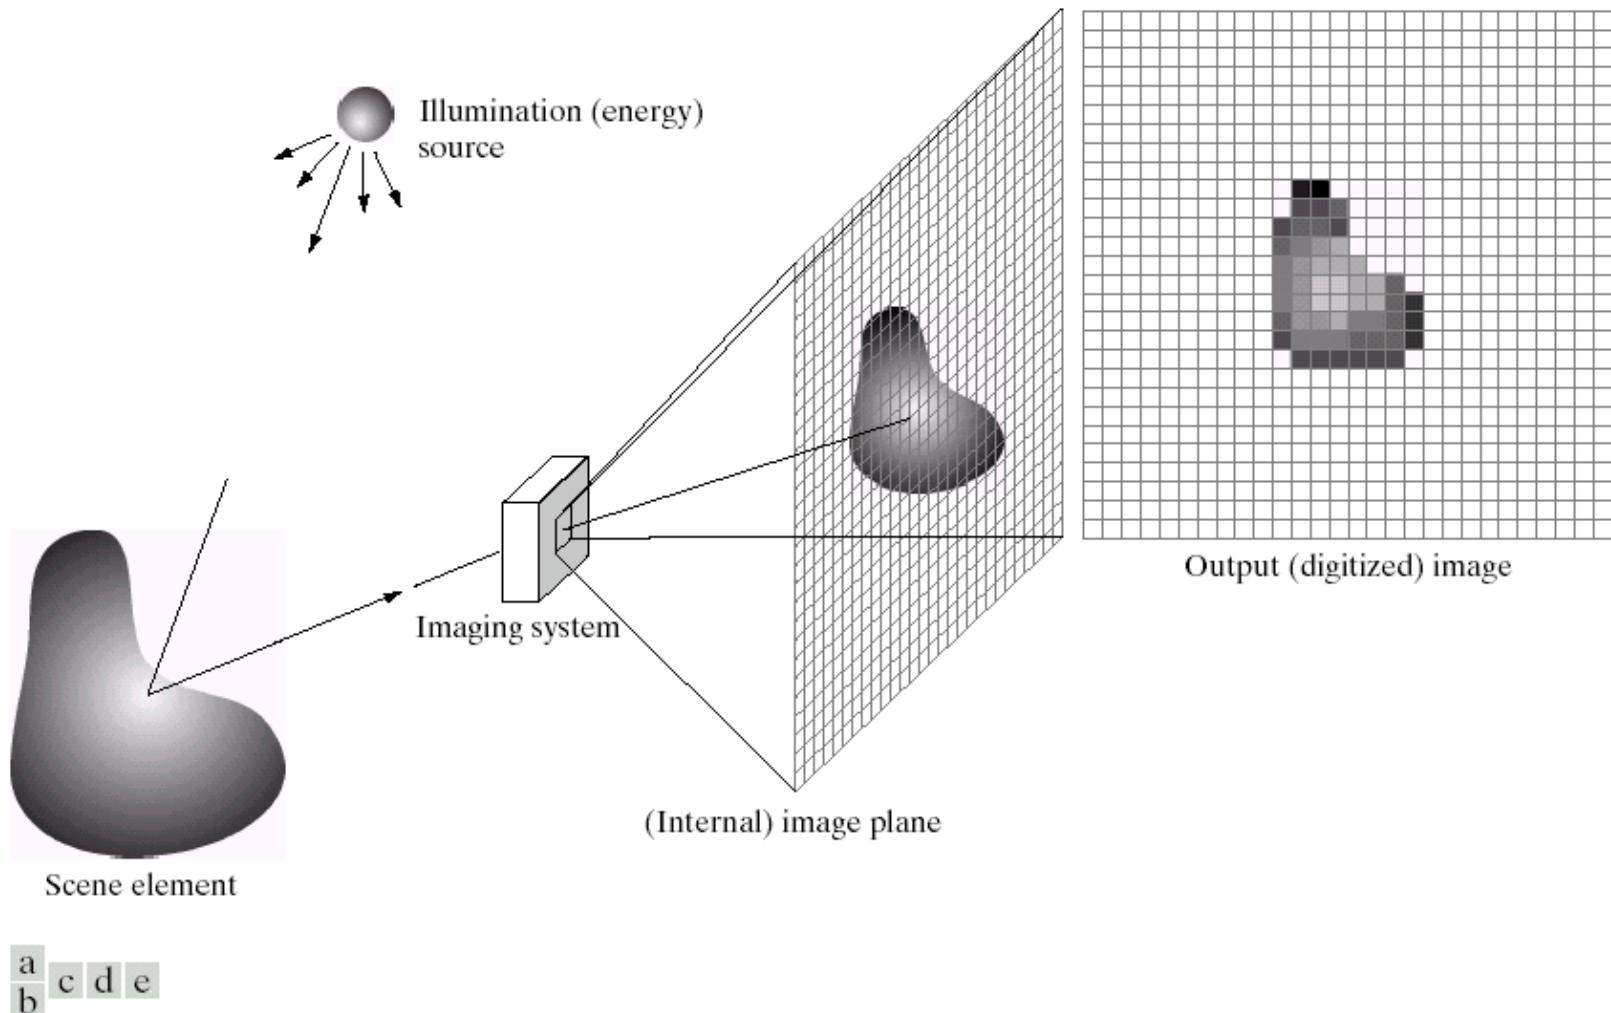
$Y$

$C_b$

$C_r$

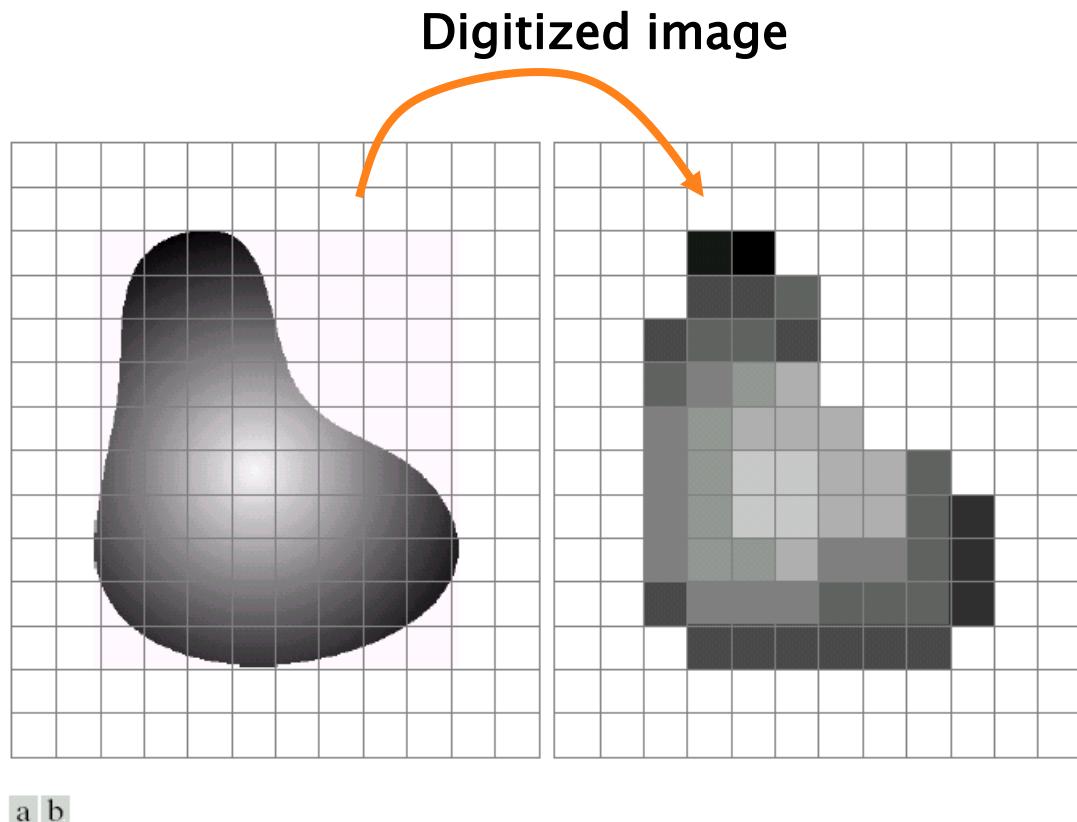
# Color Quantization

# Image formation



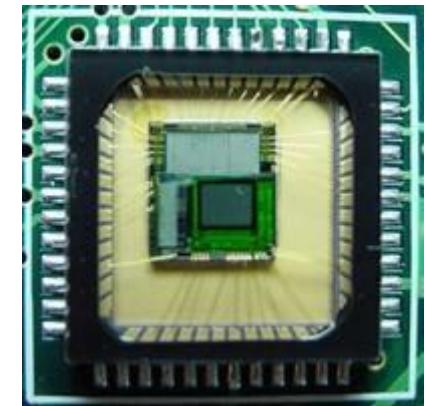
**FIGURE 2.15** An example of the digital image acquisition process. (a) Energy (“illumination”) source. (b) An element of a scene. (c) Imaging system. (d) Projection of the scene onto the image plane. (e) Digitized image.

# Sensor Array



a b

**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



CMOS sensor

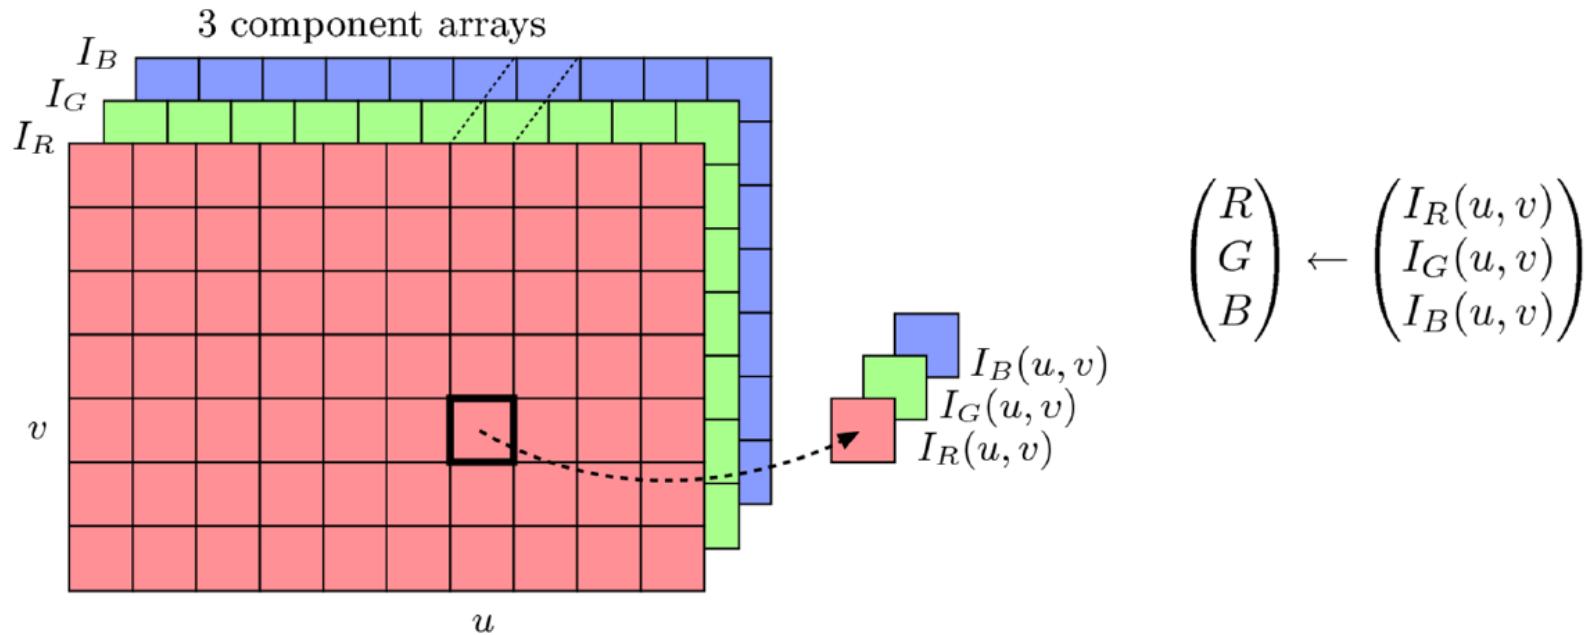
# Color Quantization

---

- ▶ **True color:** Uses all colors in color space
- ▶ **Indexed color:** Uses only some colors
  - Which subset of colors to use? Depends on application
- ▶ True color:
  - used in applications that contain many colors with subtle differences
  - E.g. digital photography or photorealistic rendering
- ▶ Two main ways to organize true color
  - Component ordering
  - Packed ordering

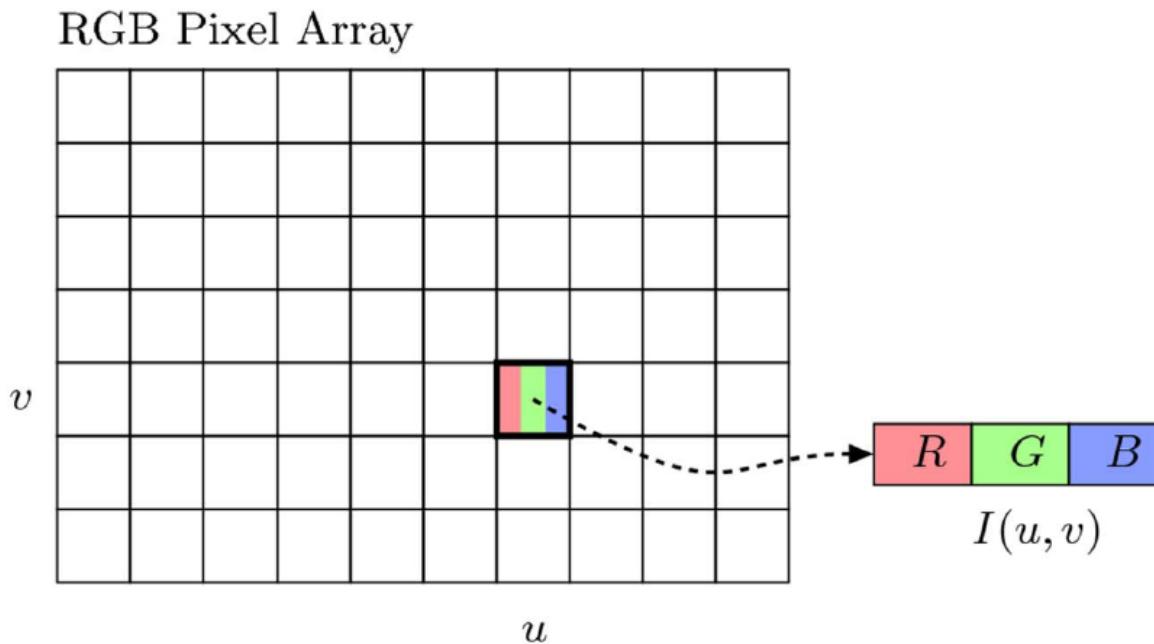
# True Color: Component Ordering

- ▶ Colors in 3 separate arrays of similar length
- ▶ Retrieve same location  $(u,v)$  in each R, G and B array



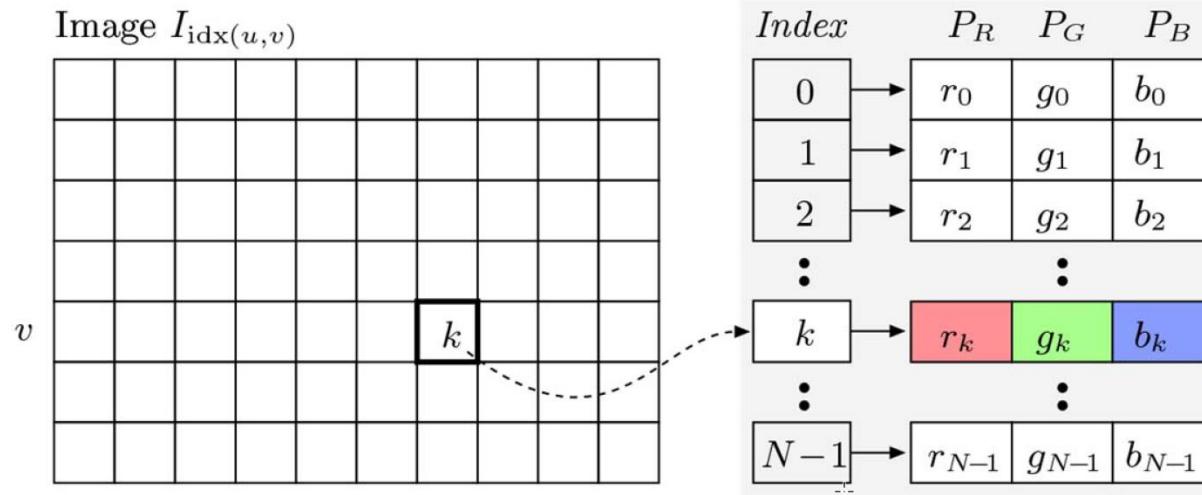
# True Color: Packed Ordering

- ▶ Component (RGB) values representing a pixel's color is packed together into single element



# Indexed color

- ▶ Permit only limited number of distinct colors ( $N = 2$  to 256)
- ▶ Used in illustrations or graphics containing large regions with same color
- ▶ Instead of intensity values, image contains indices into color table or palette
- ▶ Palette saved as part of image
- ▶ Converting from true color to indexed color requires quantization



# 1-bit Images

---

- ▶ Each pixel is stored as a single bit (0 or 1), so also referred to as **binary image**.
- ▶ Such an image is also called a 1-bit **monochrome** image since it contains no color.
- ▶ Fig. 3.1 shows a 1-bit monochrome image (called “Lena” by multimedia scientists – this is a standard image used to illustrate many algorithms).



Original image



Monochrome 1-bit Lena image.

Fig. 3.1

# 8-bit Gray-level Images

---

- ▶ Each pixel has a gray-value between 0 and 255.
  - Represented by a single byte
- ▶ **Bitmap:** The two-dimensional array of pixel values that represents the graphics/image data.
- ▶ **Image resolution** refers to the number of pixels in a digital image.
- ▶ **Frame buffer:** Hardware used to store bitmap.
  - **Video card** (actually a *graphics card*) is used for this purpose.
  - The resolution of the video card does not have to match the desired resolution of the image.



# 8-bit Gray-level Images

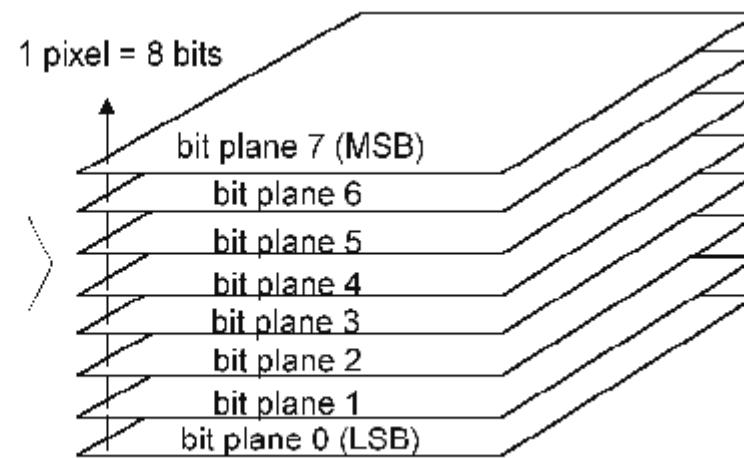
- ▶ **bit-planes:**

- Consider the 8-bit image as a set of 1-bit bitplanes
- Each plane consists of a 1-bit representation of the image

- ▶ Graphical concept of bit-planes



1 pixel = 256 grey level



Bit-planes for 8-bit grayscale image.

# Multimedia Presentation

---

- ▶ Each pixel is usually stored as a byte (a value between 0 to 255), so a  $640 \times 480$  grayscale image requires 300 kB of storage ( $640 \times 480 = 307,200$ ).
- ▶ **Dithering** : trades intensity resolution for spatial resolution to provide ability to print multi-level images on 2-level (1-bit) printers.



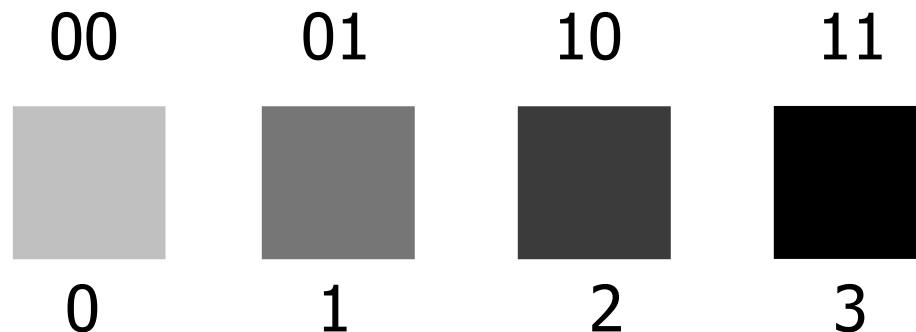
Fig. 3.3: Grayscale image of Lena.



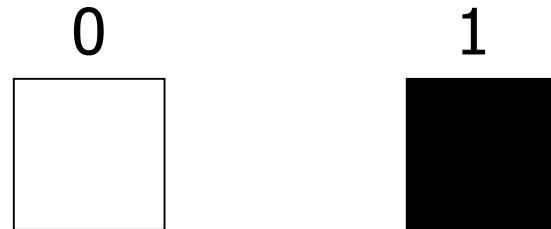
# Dither algorithm

---

- 2-bit image with 4 gray levels.

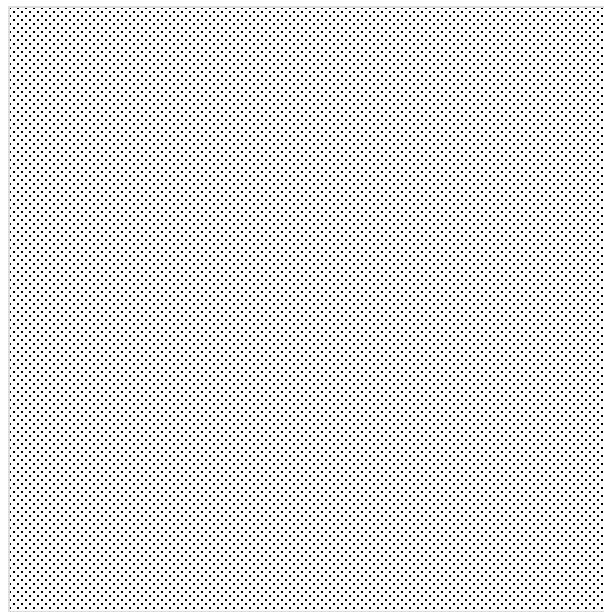
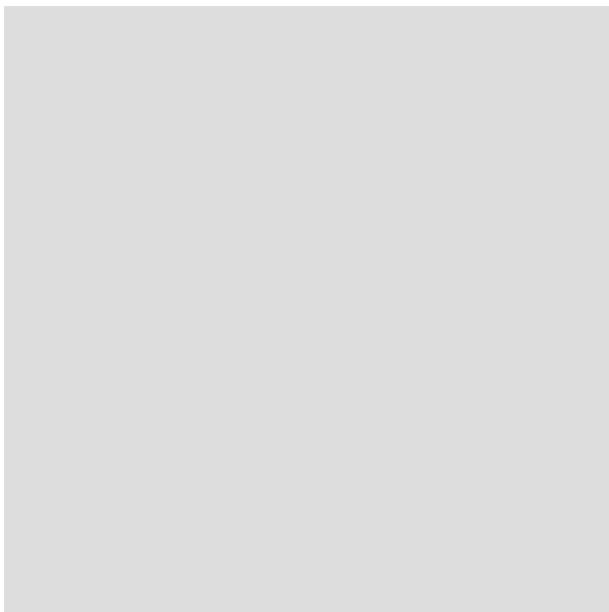


- 1-bit printer with 2 gray levels.



# Dither algorithm

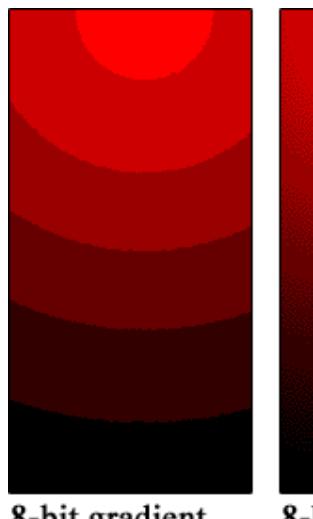
---



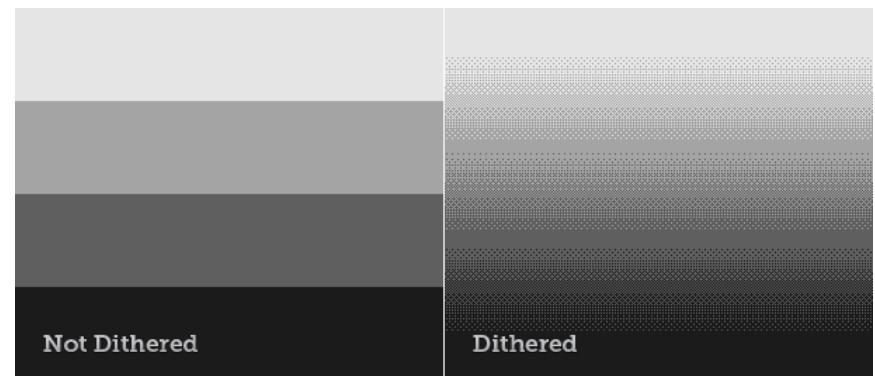
# Dither algorithm

---

## ▶ Example of dithering



Display in dark grey levels



Display using low-bit grey levels

# Dithering

---

- ▶ **Dithering** is used to calculate patterns of dots such that values from 0 to 255 correspond to patterns that are more and more filled at darker pixel values, for printing on a 1-bit printer.
- ▶ The main strategy is to replace a pixel value by a larger pattern, say 2x2 or 4x4, such that the number of printed dots approximates the varying-sized disks of ink used in analog, in **halftone printing** (e.g., for newspaper photos).
  1. Half-tone printing is an analog process that uses smaller or larger filled circles of black ink to represent shading, for newspaper printing.
  2. For example, if we use a 2x2 dither matrix

$$\begin{pmatrix} 0 & 2 \\ 3 & 1 \end{pmatrix}$$

# Ordered dither algorithm

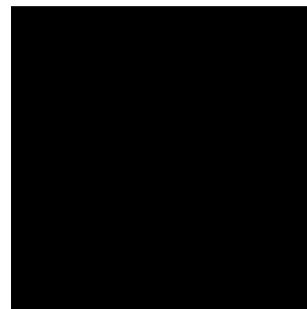
---

- Print 2-bit image using 1-bit printer based on the dither matrix

Dither Matrix

$$\begin{pmatrix} 0 & 2 \\ 3 & 1 \end{pmatrix}$$

Image



Dither result



# Ordered dither algorithm

---

- ▶ We can first re-map image values in 0..255 into the new range 0..4 by (integer) dividing by  $256/5$ . Then, e.g., if the pixel value is 0 we print nothing, in a  $2\times 2$  area of printer output. But if the pixel value is 4 we print all four dots.
- ▶ **The rule** is:
  - If (intensity > the dither matrix entry) print an **on** dot at that entry location: replace each pixel by an  $n \times n$  matrix of dots.
- ▶ Note that the image size may be much larger, for a dithered image, since replacing each pixel by a  $4 \times 4$  array of dots, makes an image 16 times as large.

# Ordered dither algorithm

---

- ▶ A clever trick can get around this problem.  
Suppose we wish to use a larger, 4x4 dither matrix,  
such as

$$\begin{pmatrix} 0 & 8 & 2 & 10 \\ 12 & 4 & 14 & 6 \\ 3 & 11 & 1 & 9 \\ 15 & 7 & 13 & 5 \end{pmatrix}$$

- ▶ An **ordered dither** consists of turning on the printer output bit for a pixel if the intensity level is greater than the particular matrix element just at that pixel position.
- ▶ Fig. 3.4 (a) shows a grayscale image of “Lena”. The ordered– dither version is shown as Fig. 3.4 (b),  
with a detail of Lena's right eye in Fig. 3.4 (c).

# Ordered dither algorithm

---

- ▶ An algorithm for ordered dither, with  $n \times n$  dither matrix, is as follows:

BEGIN

```
for  $x = 0$  to  $x_{max}$           // columns
    for  $y = 0$  to  $y_{max}$       // rows
         $i = x \bmod n$ 
         $j = y \bmod n$ 
        //  $I(x, y)/16$  is the input,  $O(x, y)$  is the output,
        // D is the dither matrix.
        if  $I(x, y) > D(i, j)$ 
             $O(x, y) = 1;$ 
        else
             $O(x, y) = 0;$ 
```

END

# Ordered dither algorithm

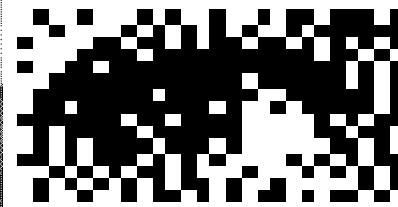
---



(a)



(b)



(c)

Fig. 3.4: Dithering of grayscale images.

(a): 8-bit grey image “lenagray.bmp”. (b): Dithered version of the image. (c): Detail of dithered version.

- ▶ Reference: <http://www.cs.indiana.edu/~dmiguse/Halftone/>  
<http://www.lenna.org>

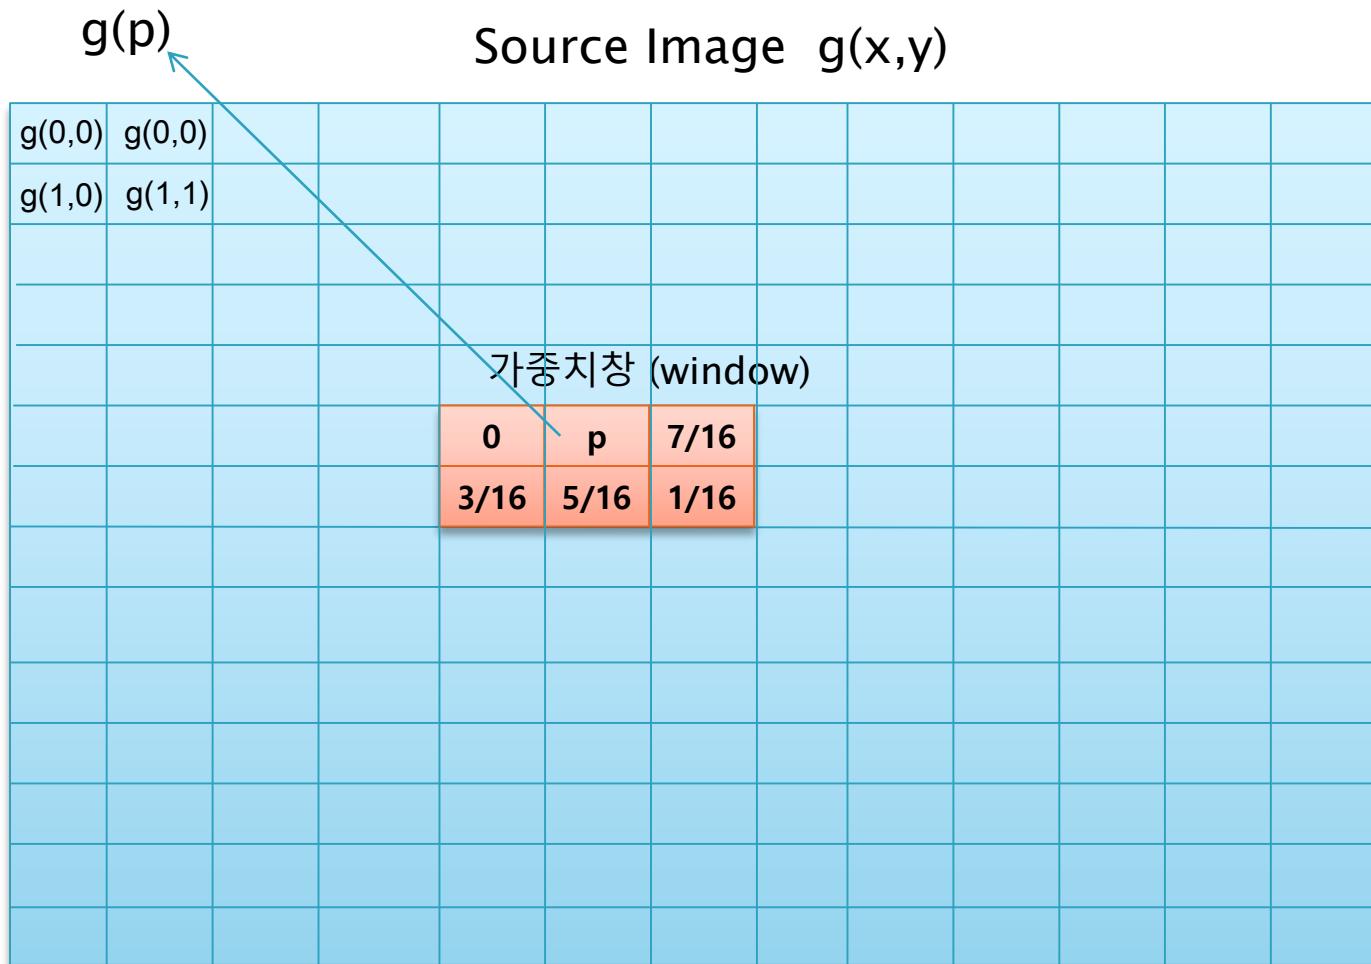
# Error diffusion dithering

- ▶ 오차확산 디더링은 픽셀값과 가까운 최대, 최소값(255 또는 0)과의 오차를 이웃한 픽셀로 확산하는 알고리듬이다. 그림에서 픽셀 p의 오차는 이웃하는 픽셀로 오차가 그림과 같은 가중치를 가지고 확산된다.

0	p	7/16
3/16	5/16	1/16

- ▶ 픽셀 p의 값이  $g(p)$ 라 할 때  $g(p)$ 가 127보다 크면  $g(p)-255$ 가 오차가 되고, 작으면  $g(p)$ 가 오차가 된다.
- ▶ 오차에 가중치를 곱하여 원영상의 위치 p의 주변 픽셀에 더해주는 과정을 모든 픽셀에 대하여 처리하면 첫 단계가 종료된다.
- ▶ 두 번째 과정은 각 픽셀의 값이 127보다 적으면 0으로 127보다 크면 1로 변환한다.

# Error diffusion dithering



# Comparison of different dither algorithms

---



Original one  
( 8-bit)



Ordered dither  
(1-bit)



Error diffusion dither  
(1-bit)

# Color Dithering (Error Diffusion)

---



Original image



Dithered image

# Color Image Types

---

- ▶ The most common data types for graphics and image file formats – **24-bit color** and **8-bit color**.
  - ▶ Some formats are restricted to particular hardware/operating system platforms, while others are “cross-platform” formats.
- 
- ▶ Even if some formats are not cross-platform, there are conversion applications that will recognize and translate formats from one system to another.
  - ▶ Most image formats incorporate some variation of a **compression** technique due to the large storage size of image files. Compression techniques can be classified into either **lossless** or **lossy**.

# 24-bit Color Images

---

- ▶ In a color 24-bit image, each pixel is represented by three bytes, usually representing RGB.
  - This format supports 256x256x256 possible colors
  - 640x480 24-bit color image requires 921.6 kB
- ▶ An important point: 24-bit color images are stored as 32-bit images, an *alpha* value representing special effect information
- ▶ Fig. 3.5 shows the image ‘forestfire.bmp’, a 24-bit image in Microsoft Windows BMP format. Also shown are the grayscale images for just the Red, Green, and Blue channels, for this image.



Fig. 3.5 High-resolution color and separate R, G, B color channel images. (a): Example of 24-bit color image “forestfire.bmp”. (b, c, d): R, G, and B color channels for this image.

# 8-bit Color Images

---

- ▶ Many systems can make use of 8 bits of color information (the so-called “256 colors”) in producing a screen image.
- ▶ Such image files use the concept of a **lookup table** to store color information.
  - Basically, the image stores not color, but instead just a set of bytes, each of which is actually an index into a table with 3-byte values that specify the color for a pixel with that lookup table index.
- ▶ Fig. 3.6 shows a 3D histogram of the RGB values of the pixels in “forestfire.bmp”.

# 8-bit Color Images

---

- ▶ Fig. 3.7 shows the resulting 8-bit image, in GIF format.



Fig. 3.7: Example of 8-bit color image.

- ▶ Note the great savings in space for 8-bit images, over 24-bit ones: a 640x480 8-bit color image only requires 300 kB of storage, compared to 921.6 kB for a color image (again, without any compression applied).

# Color Quantization

---

- ▶ What is color quantization?
  - A process uses a limited number of colors to represent an image.
  - The set of limited color is called a color palette.
- ▶ Application
  - displaying images with low--end display units
  - playing video clips with low--end display units
  - compressing images & video clips

# Color Look-up Tables (LUTs)

- ▶ The idea used in 8-bit color images is to store only the index, or code values, for each pixel. Then, e.g., if a pixel stores the value 25, the meaning is to go to row 25 in a color look-up table (LUT).

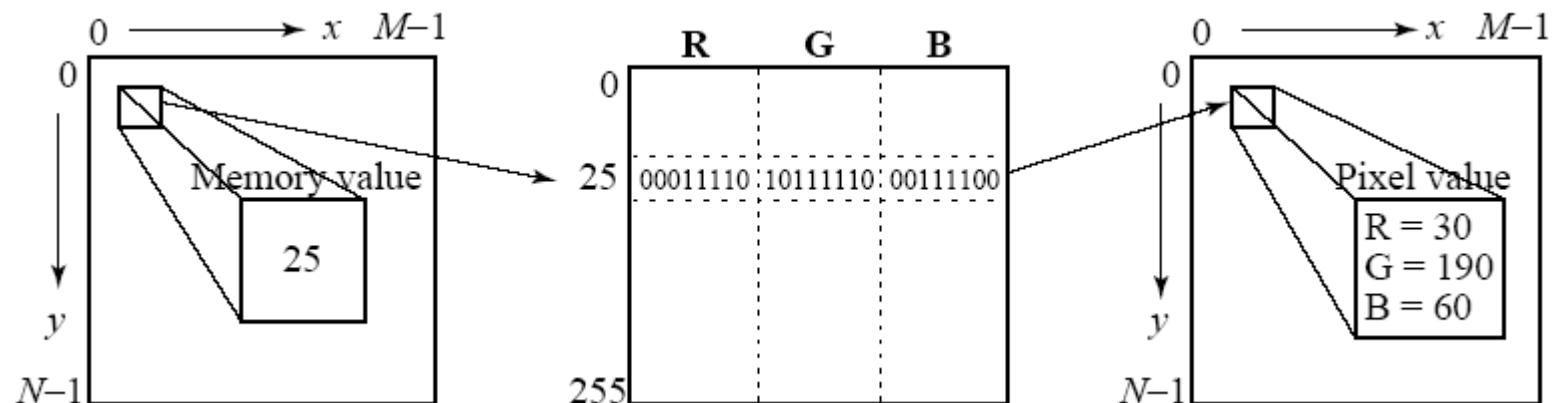
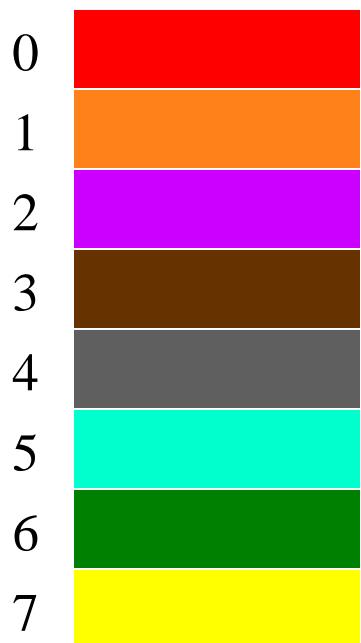


Fig. 3.8: Color LUT for 8-bit color images.

# Indexed Color

---

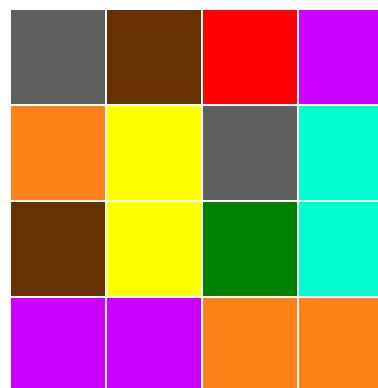
Color Table



Pixel Data

4	3	0	2
1	7	4	5
3	7	6	5
2	2	1	1

Image

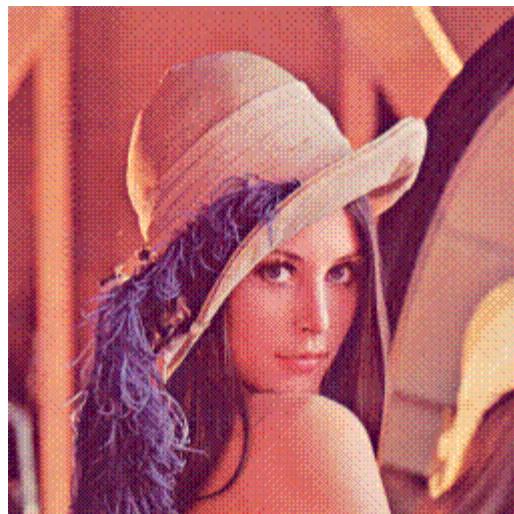


Only makes sense if you have lots  
of pixels and not many colors

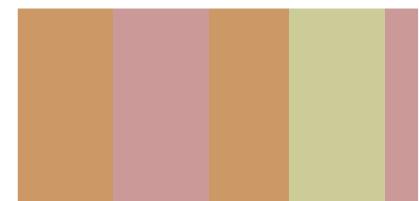
- ▶ Fig. 3.10 (a) shows a 24-bit color image of “Lena”, and Fig. 3.10 (b) shows the same image reduced to only 5 bits via dithering. A detail of the left eye is shown in Fig. 3.10 (c).



(a)



(b)



(c)

Fig. 3.10: (a): 24-bit color image “lena.bmp”. (b): Version with color dithering. (c): Detail of dithered version.

# Median-cut algorithm

---

- ▶ A simple alternate solution that does a better job for color reduction problem.
  - a. The idea is to sort the R byte values and find their median; then values smaller than the median are labeled with a “0” bit and values larger than the median are labeled with a “1” bit.
  - b. This type of scheme will indeed concentrate bits where they most need to differentiate between high populations of close colors.
  - c. One can most easily visualize finding the median by using a histogram showing counts at position 0..255.
  - d. Fig. 3.11 shows a histogram of the R byte values for the `forestfire.bmp` image along with the median of these values, shown as a vertical line.

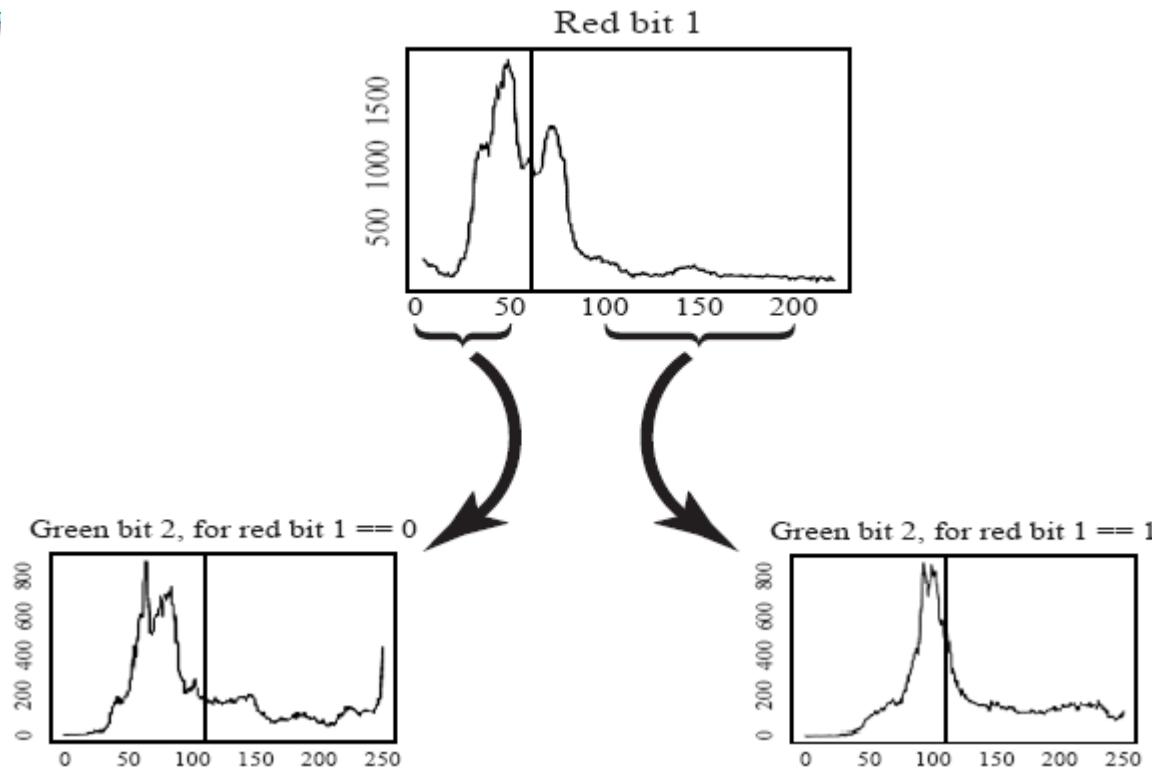
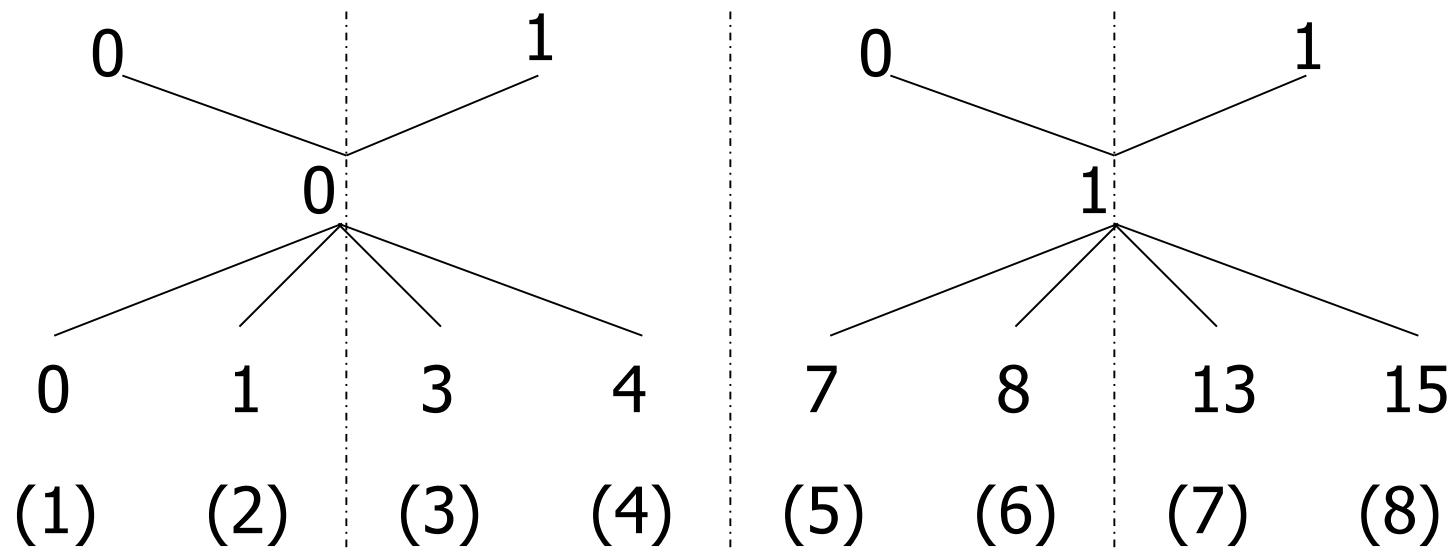


Fig. 3.11: Histogram of R bytes for the 24-bit color image “forestfire.bmp” results in a “0” bit or “1” label for every pixel. For the second bit of the color table index being built, we take R values less than the R median and label just those pixels as “0” or “1” according as their G value is less than or greater than the median of the G value, just for the “0” Red bit pixels. Continuing over R, G, B for 8 bits gives a color LUT 8-bit index.

# Median-cut algorithm

- Example (4-bit  $\rightarrow$  2-bit quantization)

3 01	1 00	4 01	0 00	15 11	8 10	13 11	7 10
---------	---------	---------	---------	----------	---------	----------	---------



# Color Quantization



2 colors



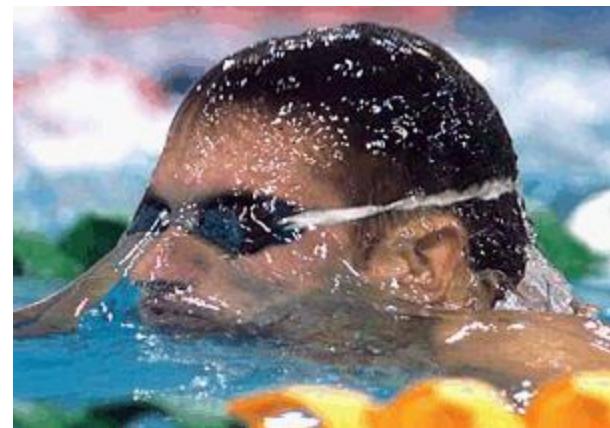
16 colors



4 colors

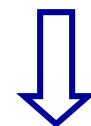


256 colors



# Quantization phases

- Sample the original image for color statistics
- Select color map based on those statistics
- Map the colors to their representative in the color map
- Redraw the image, quantizing each pixel



Algorithm



Mapping...





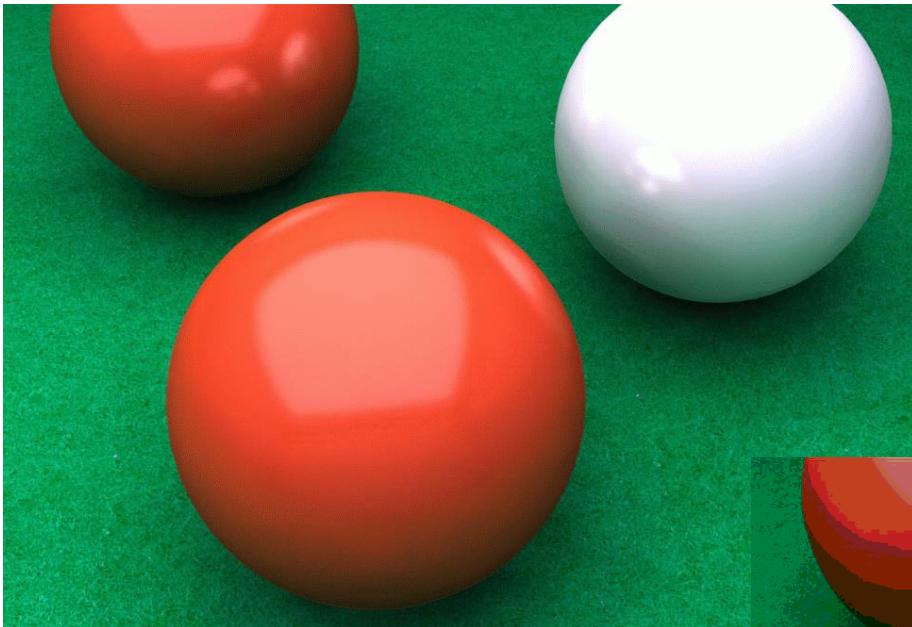
Naïve Color Quantization

24 bit to 8 bit:

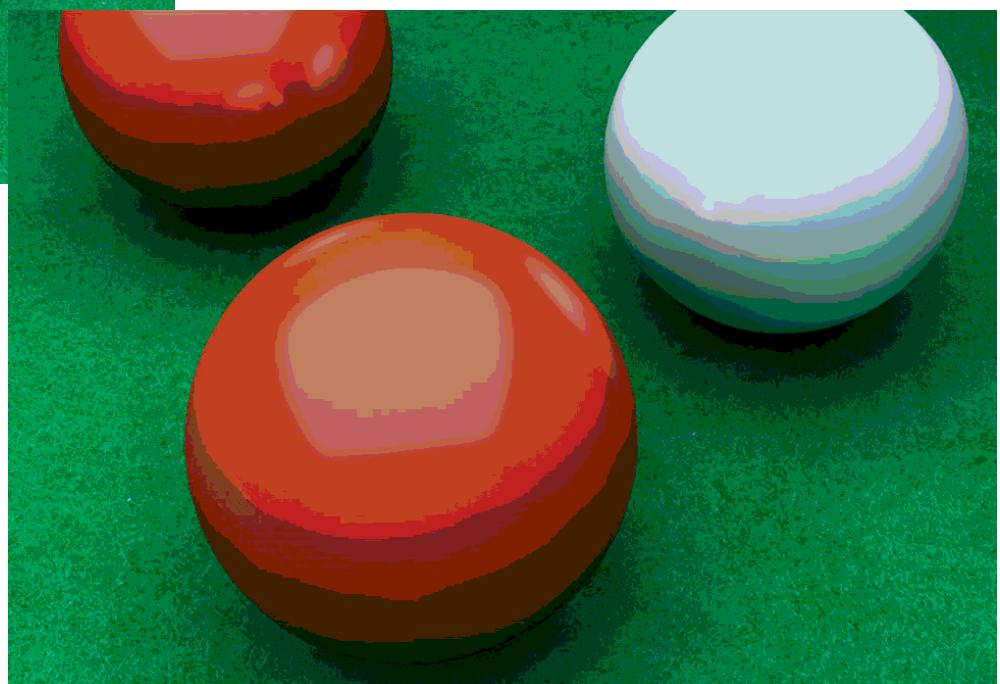
Retaining 3-3-2 most significant bits of the R, G and B components.



# 24 bit original



## 3-3-2 (fixed quantization)



24 bit original



3-3-2 (fixed  
quantization)



24 bit



Median-cut

8 bit



4 bit

