

Lecture 4

Image Processing


Multimedia Systems
Spring 2020

Digital Images

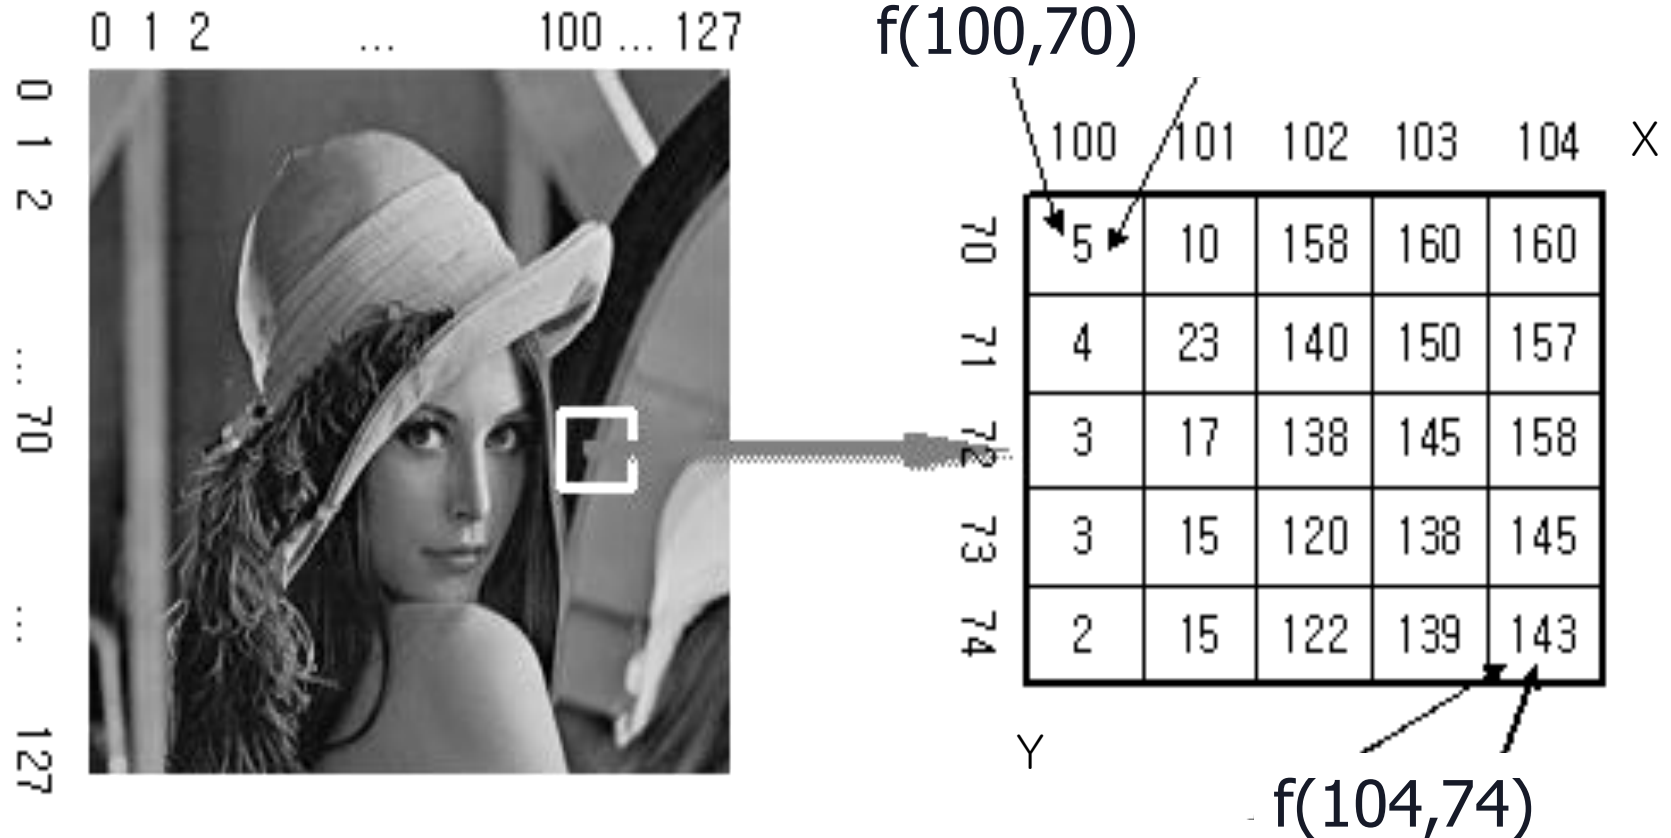
- ▶ 2D matrix of *pixels*.
 - A pixel (Picture Element) is the smallest unit of color in a digital image.
 - Number of pixels determines the *resolution*.
- ▶ Grey Scale Images:
 - A single value corresponding to a grey level.
 - 8 bits is usually used to represent a grey scale
- ▶ Color Images
 - Three values corresponding to some RGB levels
 - Red Green and Blue (RGB)
 - 24 bits to represent a color

Coordinates in an image

- ▶ Image representation
→ 2D function
- ▶ $f(x,y)$
 - $f(0,0)$: 1st col. 1st row
 - $f(1,0)$: 2nd col. 1st row
 - $f(0,1)$: 1st col. 2nd row




Coordinates in an image



Point and Neighborhood

- ▶ 3x3 neighborhood about a point (x,y) in an image point $f(x,y)$
- ▶ 8 neighborhood
 $f(x-1,y-1)$, $f(x,y-1)$, $f(x+1,y-1)$,
 $f(x-1,y)$, $f(x+1,y)$,
 $f(x-1,y+1)$,
 $f(x,y+1)$,
 $f(x+1,y+1)$



Point Operation (Binarization)

- ▶ Point operation: $p(x,y) = \text{function}(q(x,y))$
- ▶ Threshold
 - $$p(x,y) = \begin{cases} 255 & \text{if } q(x,y) > \text{threshold} \\ 0 & \text{else} \end{cases}$$




Original

threshold

Point Operation

- ▶ Image inversion
 - $p(x,y) = 255 - q(x,y)$



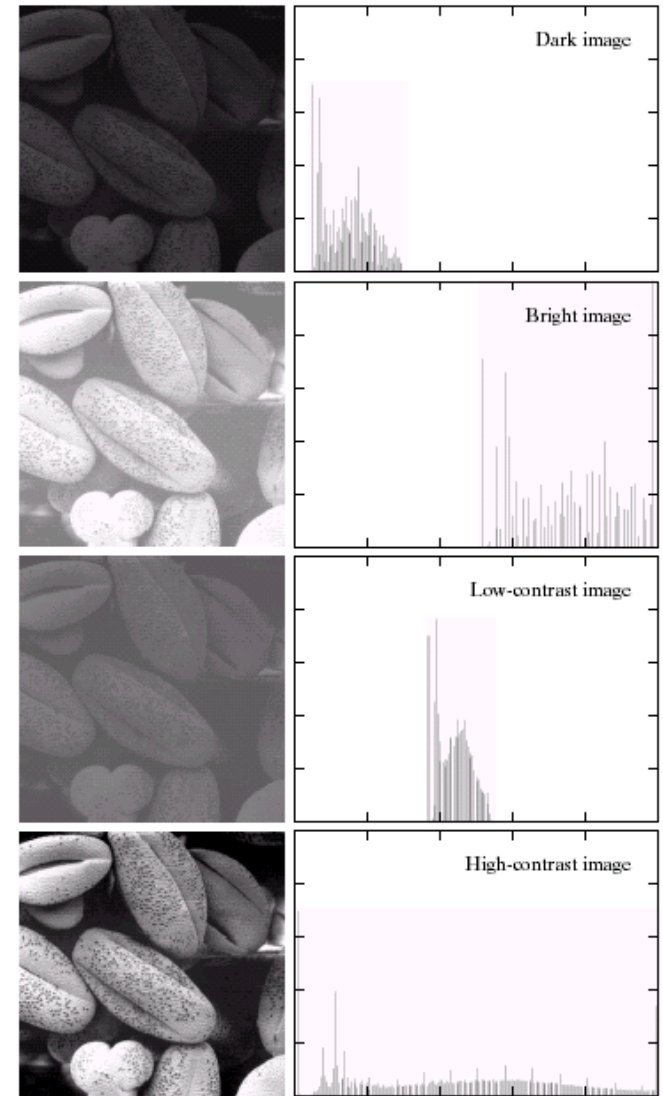

Original

Inverted

Point Operation


► Histogram Stretching

- $p(x,y) = a_1 q(x,y)$ $0 < q(x,y) \leq 100$
- $p(x,y) = a_2 q(x,y) + 100$ $100 < q(x,y) \leq 200$
- $p(x,y) = a_3 q(x,y)$ $200 < q(x,y) \leq 255$



Area operation: Image Convolution


Input Image



$$\star \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} =$$

$h(i,j)$

Output Image



Area operation: Image Convolution

- ▶ Idea: new pixel values of the image are determined by considering the surrounding pixel values.

$$p(x, y) = q * h = \sum_{i=-\frac{m}{2}}^{\frac{m}{2}} \sum_{j=-\frac{m}{2}}^{\frac{m}{2}} h(j, i) q(x + j, y + i) \quad \text{correlation}$$

$$p(x, y) = q * h = \sum_{i=-\frac{m}{2}}^{\frac{m}{2}} \sum_{j=-\frac{m}{2}}^{\frac{m}{2}} h(j, i) q(x - j, y - i) \quad \text{convolution}$$


where h is the filter, $*$ denotes convolution, m is the filter size, and $m/2$ is integer division

Linear filtering



original

q




h

?

Linear filtering



original




Filtered
(no change)

Linear filtering



original





?

shift



original




shifted

Linear filtering



original




?

Blurring






original




Blurred (filter applied in both dimensions).

Blur examples




Blur examples


impulse



original




2.4





filtered

edge



original




filtered

Linear filtering




original



h_1

—




h_2

?


Linear filtering (no change)



original



—




Filtered
(no change)


Linear filtering



original



—




?

(remember blurring)



original




Blurred (filter applied in both dimensions).


Sharpening



original




—





Sharpened
original

Sharpening example




original




Sharpened
(differences are
accentuated; constant
areas are left untouched).

Sharpening



before



after

Spatial filtering

- ▶ Representation of a general 3x3 spatial filter mask

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9


Image Convolution

► Average Smoothing:

$$\circ A_{\text{avg}} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

► Gaussian Smoothing:

$$\circ A_{\text{Gaus}} = \frac{1}{106} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 9 & 9 & 9 & 1 \\ 1 & 9 & 18 & 9 & 1 \\ 1 & 9 & 9 & 9 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$




Gaussian filter

- ▶ Smoothing (lowpass) filter
- ▶ The Gaussian smoothing operator is a 2-D convolution operator that is used to 'blur' images and remove detail and noise
- ▶ 1D Gaussian example (sigma=1.0)

0.005	0.054	0.242	0.398	0.242	0.054	0.005
-------	-------	-------	-------	-------	-------	-------

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



2D Gaussian filter (sigm1=1.0)

Gaussian filter




Image Convolution Example

- ▶ Average Smoothing with 3x3 filter:

$q(x,y)$		$h(x,y)$		$p(x,y)$																																																	
<table><tr><td></td><td>2</td><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>25</td><td>1</td><td>1</td><td></td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td></td></tr><tr><td>2</td><td>1</td><td>2</td><td>1</td><td></td></tr></table>		2	1	2	1	0	25	1	1		1	2	2	1		2	1	2	1		\ast	<table><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	$=$	<table><tr><td></td><td>2</td><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>4</td><td>4</td><td>1</td><td></td></tr><tr><td>1</td><td>4</td><td>4</td><td>1</td><td></td></tr><tr><td>2</td><td>1</td><td>2</td><td>1</td><td></td></tr></table>		2	1	2	1	0	4	4	1		1	4	4	1		2	1	2	1	
	2	1	2	1																																																	
0	25	1	1																																																		
1	2	2	1																																																		
2	1	2	1																																																		
1	1	1																																																			
1	1	1																																																			
1	1	1																																																			
	2	1	2	1																																																	
0	4	4	1																																																		
1	4	4	1																																																		
2	1	2	1																																																		

$$p(x, y) = q \ast h = \sum_{i=-\frac{m}{2}}^{\frac{m}{2}} \sum_{j=-\frac{m}{2}}^{\frac{m}{2}} h(i, j) q(x-i, y-j)$$

[illegible]

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

[illegible]

Image Convolution Example



Gaussian blur filter effect

Sharpening

$$\text{mask} \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix} \text{ or } \begin{pmatrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{pmatrix} \text{ or } \begin{pmatrix} -k & -k & -k \\ -k & 8k+1 & -k \\ -k & -k & -k \end{pmatrix}$$



Edge Detection

- ▶ Derivative of a linear function $f(x,y)$

$$\frac{\partial f}{\partial x} = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon} \right)$$

- ▶ We could approximate this as

$$\frac{\partial f}{\partial x} \approx \frac{f(x_{n+1}, y) - f(x_n, y)}{\Delta x}$$

- ▶ In image, this can be
 - $I(x+1, y) - I(x, y)$ or $I(x+1, y) - I(x-1, y)$

Simple Edge Detection

▶ Simple Edge Detection Idea:

1. $\partial I / \partial x$

Simply convolve Image with $[1 \ 0 \ -1]$

Resulting image: I_x

2. $\partial I / \partial y$

Simply convolve Image with $[1 \ 0 \ -1]^T$

Resulting image: I_y

3. ***Image Gradient (magnitude $\partial I / \partial x$ and $\partial I / \partial y$) :***

$$\nabla I(i,j) = \sqrt{I_x(i,j)^2 + I_y(i,j)^2}$$


$$\partial \mathbf{I} / \partial \mathbf{x}$$

	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	-9	-9	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	

$$\partial \mathbf{I} / \partial \mathbf{y}$$

	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	9	9	9	9	9	9	
	0	0	0	0	9	9	9	9	9	9	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	

$$\nabla \mathbf{I}$$



	0	0	0	0	9	9	0	0	0	0	
	0	0	0	0	9	9	0	0	0	0	
	0	0	0	0	9	9	0	0	0	0	
	0	0	0	0	9	9	0	0	0	0	
	0	0	0	0	9	9	0	0	0	0	
	0	0	0	0	9	9√2	9	9	9	9	
	0	0	0	0	0	9	9	9	9	9	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	

$$\sqrt{I_x(i,j)^2 + I_y(i,j)^2}$$



∇I



Sobel Edge Detection

- ▶ Sobel edge filter

$$[\mathbf{h}_x] = \frac{1}{4} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \cdot [1 \ 0 \ -1]$$

$$[\mathbf{h}_y] = \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \frac{1}{4} \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} \cdot [1 \ 2 \ 1]$$


$$p(x, y) = \sqrt{(q(x, y) * h_x(x, y))^2 + (q(x, y) * h_y(x, y))^2}$$

$$\theta = \arctan\left(\frac{q(x, y) * h_x(x, y)}{q(x, y) * h_y(x, y)}\right)$$




Laplacian Filters


- ▶ Let us take a look at the one-dimensional case:
- ▶ A change in brightness:



- Its first derivative:



- Its second derivative:



Laplacian Filters

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

$$\frac{\partial^2 f}{\partial x^2} = \frac{\partial G_x}{\partial x} = \frac{\partial(f[i, j+1] - f[i, j])}{\partial x}$$

$$= \frac{\partial(f[i, j+1])}{\partial x} - \frac{\partial f[i, j]}{\partial x}$$

$$= (f[i, j+2] - f[i, j+1]) - (f[i, j+1] - f[i, j])$$

$$= f[i, j+2] - 2f[i, j+1] + f[i, j]$$

Laplacian Filters

- ▶ Since we want the computation centered at $[i,j]$:


$$\frac{\partial^2 f}{\partial x^2} = f[i, j+1] - 2f[i, j] + f[i, j-1]$$

$$\frac{\partial^2 f}{\partial y^2} = f[i+1, j] - 2f[i, j] + f[i-1, j]$$

- Obviously, each function can be computed using a 1, -2, 1 convolution filter in x or y direction, respectively.
- Since ∇^2 is defined as the sum of these two functions, we can use a single 3×3 (or larger) convolution filter for its computation (next slide).

Laplacian Filters

► Continuous variant:



Discrete variants
(applied after smoothing):

0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1

Laplacian of Gaussian

- ▶ Bad idea to apply a Laplacian without smoothing
 - smooth with Gaussian, apply Laplacian
 - this is the same as filtering with a Laplacian of Gaussian (LOG) filter
 - Now mark the zero points *where there is a sufficiently large derivative, and enough contrast*



Laplacian of Gaussian




▶ 3×3 Laplacian



■ 5×5 Laplacian




■ 7×7 Laplacian



■ zero detection




■ zero detection



■ zero detection

Median filter


- ▶ A **Median Filter** operates over a window by selecting the median intensity in the window.



- ▶ What advantage does a median filter have over a mean filter?
- ▶ Is a median filter a kind of convolution?
- ▶ Median filter is non linear

Median filter

- ▶ Useful in removing salt-and-pepper noise
- ▶ But, salt-and-pepper noise is not common.
- ▶ Also useful in removing noise in a binary image



Template(Block) matching

Template Matching

- ▶ Filters can be regarded as templates
 - Applying a filter at some point can be seen as taking a dot-product between the image and some vector
 - Filtering the image is a set of dot products
 - It is a measure of the angle of two vectors (template (block) and image)

Template Matching by Cross-Correlation

- ▶ The use of cross-correlation for template matching is motivated by the distance measure (squared Euclidean distance) between a template $t(x,y)$ and image $f(x,y)$

$$d_{f,t}(x, y) = \sum_{i,j} [t(i, j) - f(x + j, y + i)]^2$$

- ▶ The correlation response at each pixel (x,y) is

$$c(x, y) = f * t = \sum_{i=-\frac{m}{2}}^{\frac{m}{2}} \sum_{j=-\frac{m}{2}}^{\frac{m}{2}} t(j, i) f(x + j, y + i)$$

Normalize correlation

- ▶ By normalizing the image and feature vectors to unit length, yielding a cosine-like correlation coefficient

$$\gamma(x, y) = \frac{\sum_{i,j} [(t(j, i) - \bar{t})(f(x + j, y + i) - \bar{f}_{x,y})]}{\{\sum_{i,j} (t(j, i) - \bar{t})^2 \sum_{i,j} (f(x + j, y + i) - \bar{f}_{x,y})^2\}^{0.5}}$$

$$-1 < \gamma(x, y) \leq 1$$


- 1 or -1 : exact matching of image and feature
- 0: no correlation

Template matching: example(1)

Template




Original image



Absolute value of correlation
(scale to 0 ~ 255)

Template matching: example(2)

Template



Original image




Absolute value of correlation
(scale to 0 ~ 255)



Template matching: example(3)



Template



Data Set 3



Normalized correlation

Correspondence by block matching

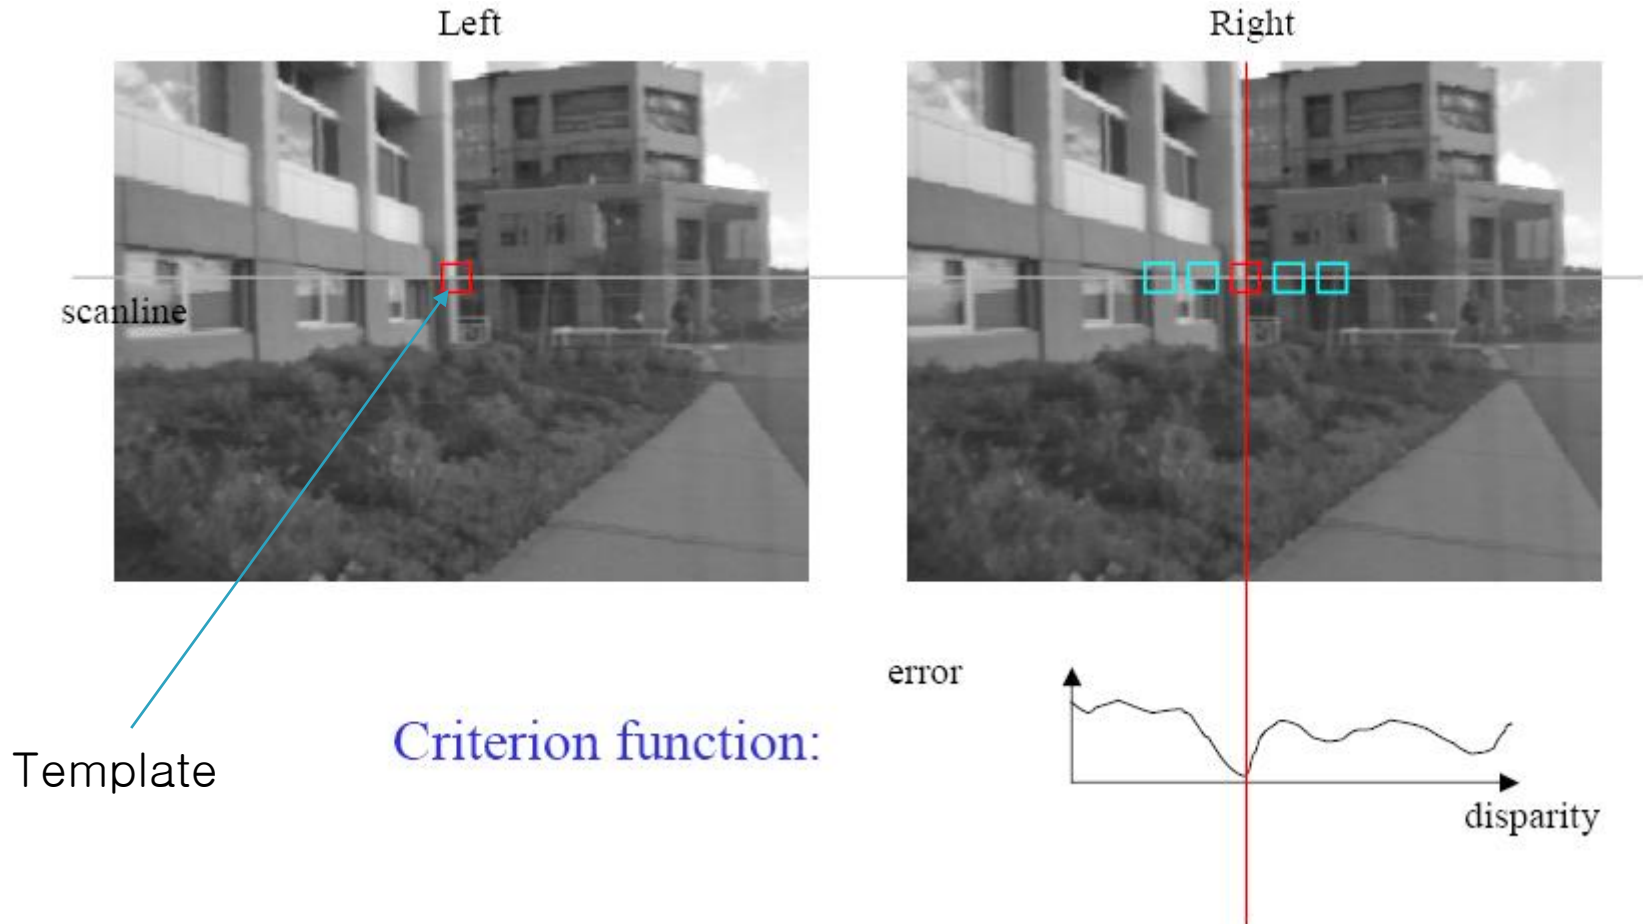


Image blocks as a vector

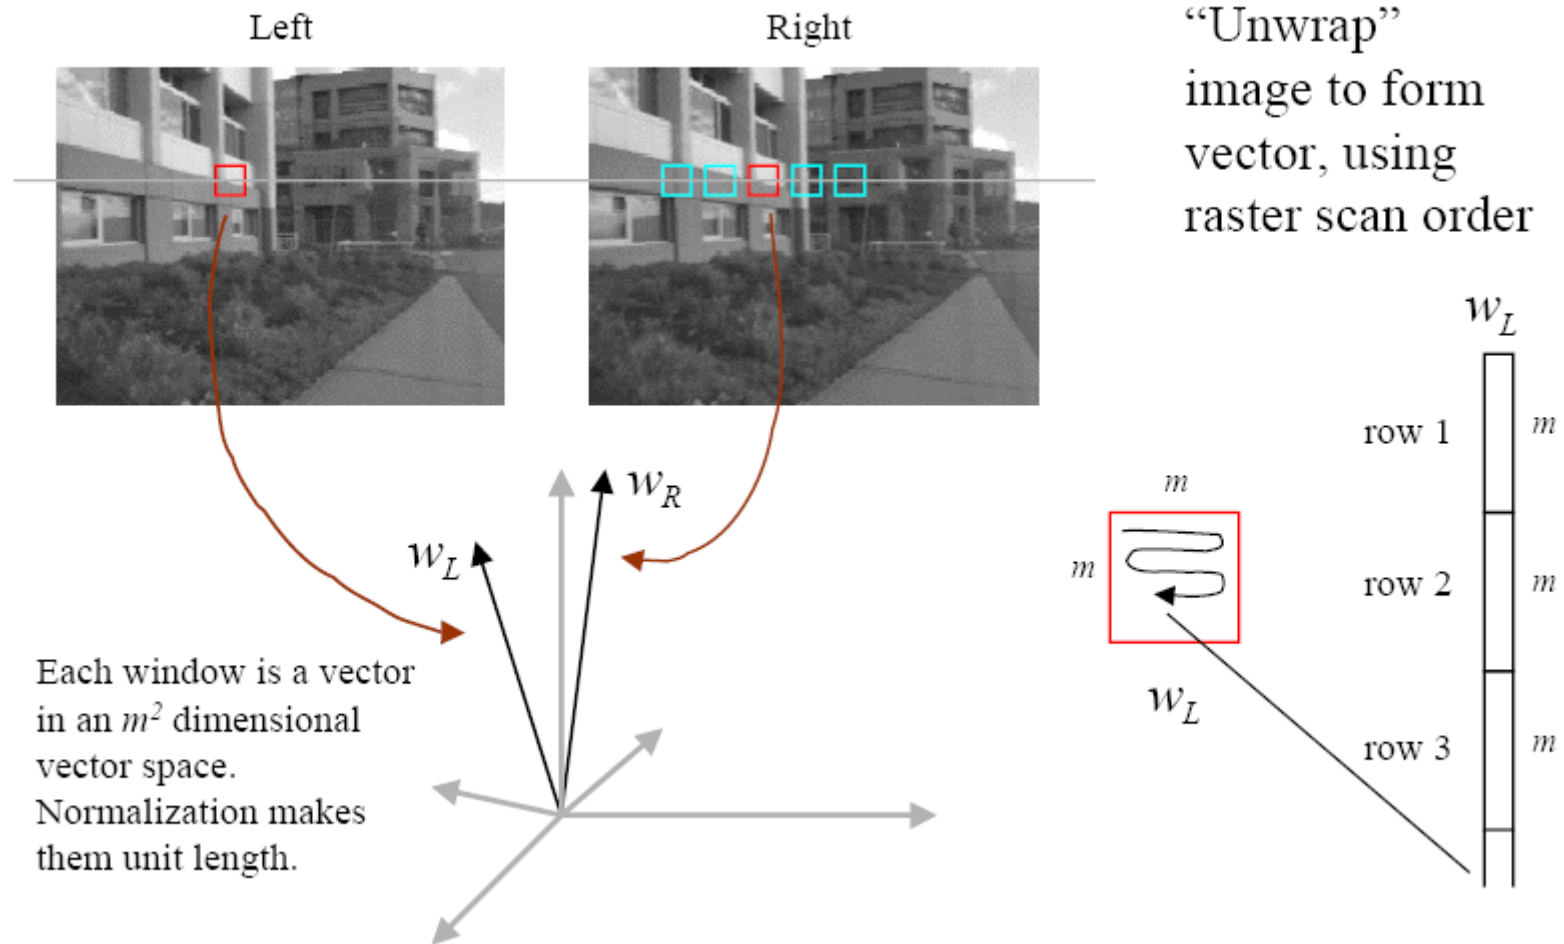




Image blocks as a vector



Sum of squared distance



w_L and w_R are corresponding m by m windows of pixels.

We define the window function :

$$W_m(x, y) = \{u, v \mid x - \frac{m}{2} \leq u \leq x + \frac{m}{2}, y - \frac{m}{2} \leq v \leq y + \frac{m}{2}\}$$

The SSD cost measures the intensity difference as a function of disparity :

$$C_r(x, y, d) = \sum_{(u, v) \in W_m(x, y)} [I_L(u, v) - I_R(u - d, v)]^2$$